

Getting Technical with Spark



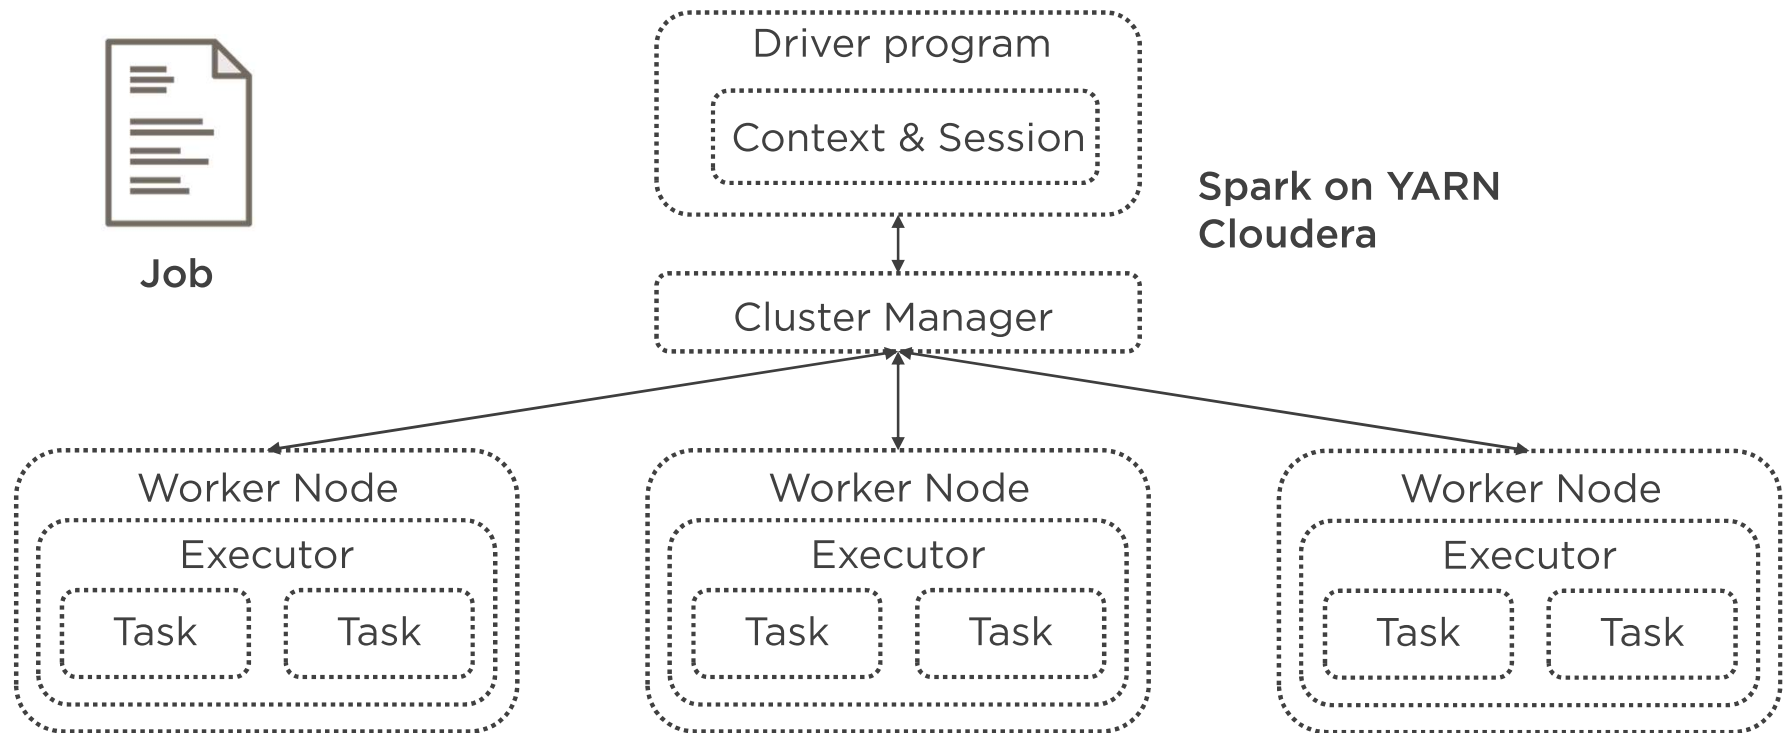
Xavier Morera

HELPING DEVELOPERS UNDERSTAND SEARCH & BIG DATA



@xmorera www.xaviermorera.com




Spark's Architecture





Developing Spark Applications



What's Missing?



Storage



Reading from Different Data Sources

```
sc.textFile("/user/cloudera/spark-committers-no-  
header.tsv").take(10)
```

```
sc.textFile("file:///stackexchange/spark-committers-no-  
header.tsv").take(10)
```

```
sc.textFile("s3a://pluralsight-spark-cloudera-scala/spark-  
committers-no-header.tsv").take(10)
```




Let's Talk APIs




Let's Talk APIs

RDD



Since the beginning

DataFrame (and Dataset)





Search for messages



Groups



Spark Users ›

glom?

5 posts by 3 authors



- ★ **Mark Hamstra** In reading through the Spark source to discover all of the undocumented goodies, I've come across the `glom()` method of RDDs. Okay, so I can turn an RDD into a `GlommedRDD`. 9/11/12
- ★ **Scott Smith** `x.glom().flatMap(_ -> _) == x` It's great for taking all the rows of a split and turning them into a single larger object for doing more efficient operations than `map` would allow (like matrix 9/12/12
- ★ **Matei Zaharia** It's meant for algorithms that want to work on an RDD's entire partition at once, such as stochastic gradient descent (a way to do gradient descent where you just merge in one point 9/12/12



Mark Hamstra

9/12/12



- ★ Is there a plan to deprecate `glom` (or any thing else) in 0.6? How much more do you anticipate the api evolving before the 0.6 release?

- show quoted text -



Matei Zaharia

9/12/12




- ★ I think we'll keep all of the APIs the same. I don't like removing these kinds of operations because some user code might depend on them and they're not that hard to maintain.

Matei

- show quoted text -


Let's Talk APIs

RDD




Since the beginning

DataFrame (and Dataset)




Let's Talk APIs

RDD



Since the beginning

DataFrame (and Dataset)



Since 2.0



Upgrading From Spark SQL 2.1 to 2.2

- Spark 2.1.1 introduced a new configuration key: `spark.sql.hive.caseSensitiveInferenceMode`. It had a default setting of `NEVER_INFER`, which kept behavior identical to 2.1.0. However, Spark 2.2.0 changes this setting's default value to `INFER_AND_SAVE` to restore compatibility with reading Hive metastore tables whose underlying file schema have mixed-case column names. With the `INFER_AND_SAVE` configuration value, on first access Spark will perform schema inference on any Hive metastore table for which it has not already saved an inferred schema. Note that schema inference can be a very time consuming operation for tables with thousands of partitions. If compatibility with mixed-case column names is not a concern, you can safely set `spark.sql.hive.caseSensitiveInferenceMode` to `NEVER_INFER` to avoid the initial overhead of schema inference. Note that with the new default `INFER_AND_SAVE` setting, the results of the schema inference are saved as a metastore key for future use. Therefore, the initial schema inference occurs only at a table's first access.

Upgrading From Spark SQL 2.0 to 2.1

- Datasource tables now store partition metadata in the Hive metastore. This means that Hive DDLs such as `ALTER TABLE PARTITION ... SET LOCATION` are now available for tables created with the Datasource API.
 - Legacy datasource tables can be migrated to this format via the `MSCK REPAIR TABLE` command. Migrating legacy tables is recommended to take advantage of Hive DDL support and improved planning performance.
 - To determine if a table has been migrated, look for the `PartitionProvider: Catalog` attribute when issuing `DESCRIBE FORMATTED` on the table.
- Changes to `INSERT OVERWRITE TABLE ... PARTITION ...` behavior for Datasource tables.
 - In prior Spark versions `INSERT OVERWRITE` overwrote the entire Datasource table, even when given a partition specification. Now only partitions matching the specification are overwritten.
 - Note that this still differs from the behavior of Hive tables, which is to overwrite only partitions overlapping with newly inserted data.

Upgrading From Spark SQL 1.6 to 2.0

- `SparkSession` is now the new entry point of Spark that replaces the old `SQLContext` and `HiveContext`. Note that the old `SQLContext` and `HiveContext` are kept for backward compatibility. A new `catalog` interface is accessible from `SparkSession` - existing API on databases and



RDD or Dataset/DataFrame?

RDD

Unstructured data

Data manipulation with lambdas

Low level transformations

Functional

Dataset/DataFrame

Structured or semi structured data

Much easier to understand

Equivalent to database table

But better

Leverage optimizations

Relational

Datasets have types






RDD

DataFrame



RDD Optimizations



**Catalyst
Optimizer**

**Project
Tungsten**



Catalyst Optimizer

Framework, part of Spark SQL

Optimize queries

Objectives

- Easily add new optimization techniques
- Enable external developers to extend it



Catalyst Optimizer

Analyzes queries

Creates an unresolved logical plan

Looks at the Catalog

Creates an optimized logical plan

Optimized physical plan

Optimized code generation



Project Tungsten

Push the boundaries of performance

Optimize for CPU and memory efficiency

Focuses on hardware where Spark runs



Project Tungsten

Tungsten row format

Manages memory explicitly

Cache aware computation




Project Tungsten

Tungsten row format

Manages memory explicitly

Cache aware computation



Project Tungsten

Tungsten row format

Manages memory explicitly

Cache aware computation

Whole-stage Code Generation



SparkContext

SparkSession



SparkContext

Located in the Spark Driver

Entry point for RDDs

The Spark application

- One SparkContext per application

Created for you in the REPL

You need to create for spark2-submit



SparkContext

```
sc
sc.version
sc.getClass
sc.appName
sc.getClass // try adding .getName
sc //hit tab, autocomplete will kick in
sc.uiWebUrl
sc.applicationId
sc.sparkUser
sc.textFile("/user/cloudera/spark-committers-no-header.tsv").take(10)
sc.parallelize(Array("Matei Zaharia", "Josh Rosen", "Holden Karau"))
```



SparkSession

Entry point to Spark SQL

Merges SQLContext and HiveContext

Access SparkContext

Can have multiple SparkSession objects

Created for you in REPL

You need to create for spark2-submit



SparkSession

```
spark
```

```
spark.getClass.getName
```

```
spark.sparkContext
```

```
spark.sql("select * from committers").show(10)
```

```
cmDF.select("Name").show(10)
```

```
spark.catalog.listDatabases().show(truncate=false)
```



SparkContext

SparkSession



Spark Configuration

Properties

Application parameters

Environment Vars


Per machine settings

Logging

In log4j.properties



Configuration Precedence



Application (in code)

Flags passed to spark2-submit/spark2-shell

spark-defaults.conf



Environment

Runtime Information

Name	Value
Java Home	/usr/java/jdk1.8.0_144/jre
Java Version	1.8.0_144 (Oracle Corporation)
Scala Version	version 2.11.8

Spark Properties

Name	Value
spark.app.id	application_1511382093107_0012
spark.app.name	PySparkShell
spark.authenticate	false
spark.driver.appUIAddress	http://10.0.2.101:4040
spark.driver.extraLibraryPath	/opt/cloudera/parcels/CDH-5.12.1-1.cdh5.12.1.p0.3/lib/hadoop/lib/native
spark.driver.host	10.0.2.101
spark.driver.port	34232
spark.dynamicAllocation.enabled	true
spark.dynamicAllocation.executorIdleTimeout	60
spark.dynamicAllocation.minExecutors	0
spark.dynamicAllocation.schedulerBacklogTimeout	1
spark.eventLog.dir	hdfs://mgt01.cloudera:8020/user/spark/spark2ApplicationHistory
spark.eventLog.enabled	true
spark.executor.extraLibraryPath	/opt/cloudera/parcels/CDH-5.12.1-1.cdh5.12.1.p0.3/lib/hadoop/lib/native
spark.executor.id	driver
spark.executorEnv.PYTHONPATH	/opt/cloudera/parcels/SPARK2-2.2.0.cloudera1-1.cdh5.12.0.p0.142354/lib/spark2/python/lib/py4j-0.10.4-src.zip:/opt/cloudera/parcels/SPARK2-2.2.0.cloudera1-1.cdh5.12.0.p0.142354/lib/spark2/python/:<CPS>/opt/cloudera/parcels/SPARK2-2.2.0.cloudera1-1.cdh5.12.0.p0.142354/lib/spark2/python/lib/py4j-0.10.4-src.zip<CPS>/opt/cloudera/parcels/SPARK2-2.2.0.cloudera1-1.cdh5.12.0.p0.142354/lib/spark2/python/lib/pyspark.zip
spark.hadoop.mapreduce.application.classpath	
spark.hadoop.yarn.application.classpath	
spark.master	yarn
spark.org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter.param.PROXY_HOSTS	dn04.cloudera


```
import org.apache.spark.{SparkContext, SparkConf}
val conf = new SparkConf()
    .setMaster("yarn")
    .setAppName("Stack Overflow Test")
    .set("spark.executor.memory", "1g")
val sc = new SparkContext(conf)
```

SparkConf with RDD API

Initialize **SparkConf** object

And pass on context creation

Object is cloned, so can't be modified



Time and Size Formats



5ms for 5 milliseconds

10m for 10 minutes


1y for 1 year

1g / 1gb for 1 gigabyte

1t / 1tb for 1 terabyte



~~Time and Size Formats~~



Imagine if you had to do single value
31622400 seconds **for** 1 year



```
spark2-submit --executor-memory 4g stackoverflowtest.scala
```

SparkConf with spark2-submit

You can pass configuration options when using **spark2-submit**

Dynamically set properties on execution

Also possible to use **--config "spark.executor.memory=4g"**



```
scala> sc.getConf.getAll  
  
res35: Array[(String, String)] =  
Array((spark.driver.host,10.0.2.104),  
(spark.eventLog.enabled,true),  
(spark.driver.extraLibraryPath,/opt/cloudera/parcels/CDH-  
5.12.1-1.cdh5.12.1.p0.3/lib/hadoop/lib/native),  
  
scala> sc.getConf.get("spark.driver.host")
```

Reviewing Configuration

And you can review the configuration with **sc.getConf.getAll**

Or one specific configuration with **sc.getConf.get**



```
import org.apache.spark.sql.SparkSession  
val spark = SparkSession.  
    builder().  
    appName("StackOverflow Test").  
    config("spark.executor.memory", "1g").  
    getOrCreate()
```

SparkConf with Dataset API

Import **SparkSession**

Builder pattern: complex object, one step at a time

Named properties like **master** and **appName**



```
val spark = SparkSession.  
  builder().  
  appName("StackOverflow Test").  
  config("spark.submit.deployMode", "").  
  getOrCreate()
```

Another Useful Configuration

Deployment mode

Specifies location of where the Spark **driver** executes



```
val spark = SparkSession.  
    builder().  
    appName("StackOverflow Test").  
    config("spark.submit.deployMode", "client").  
    getOrCreate()
```

Client Deployment Mode

Driver runs where Spark application was launched

- i.e. local machine

Process killed if **driver** disconnected




```
val spark = SparkSession.  
    builder().  
    appName("StackOverflow Test").  
    config("spark.submit.deployMode", "cluster").  
    getOrCreate()
```

Cluster Deployment Mode

Driver runs in a node within the cluster even if launched from outside

Process not killed if computer where submitted is disconnected

Does not support shell



Spark Configuration

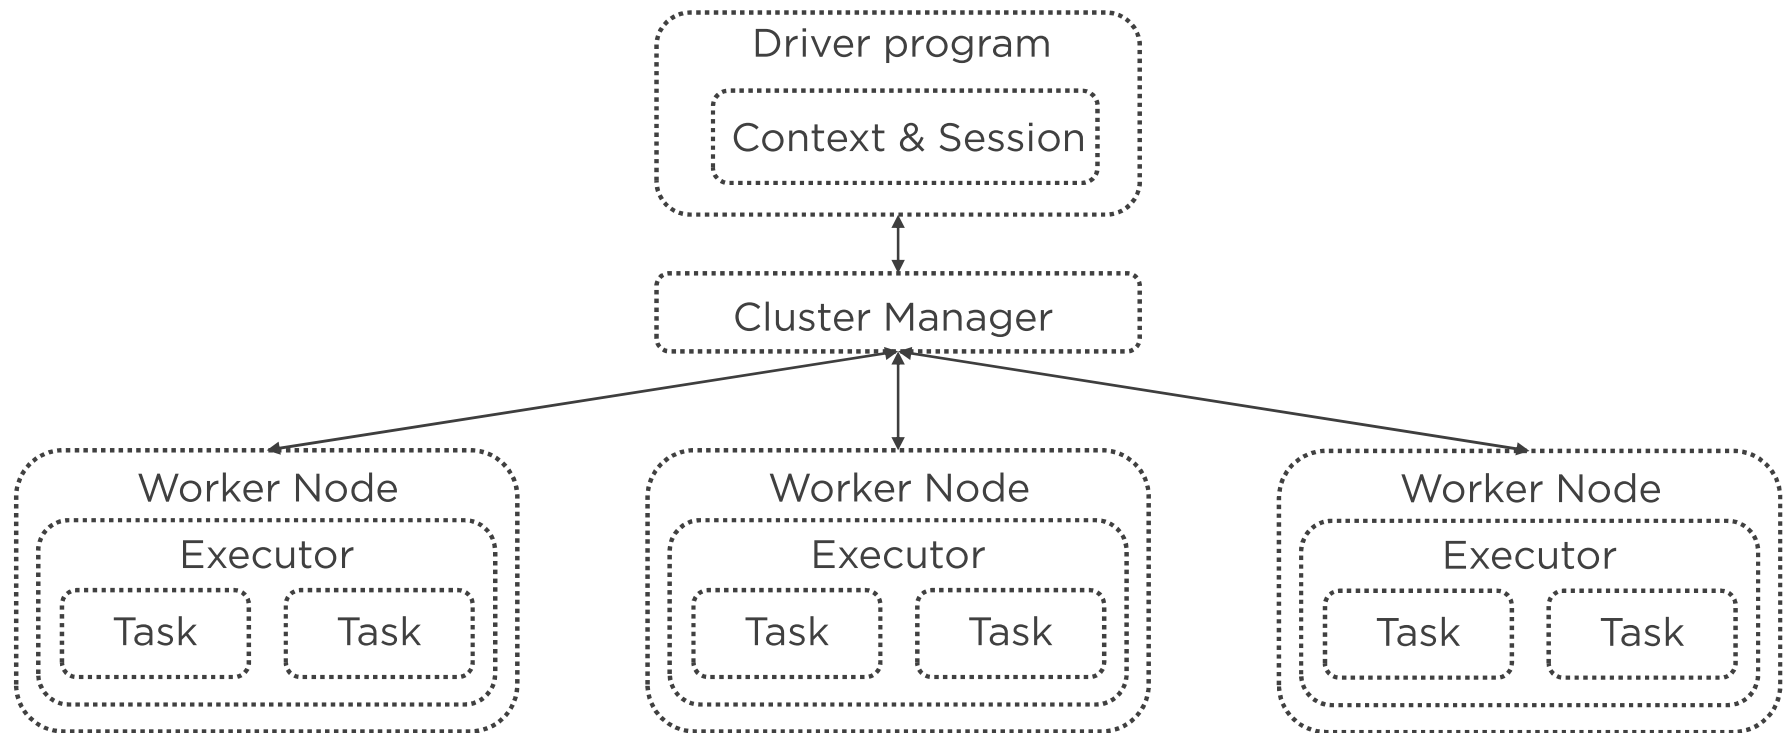


Configure Spark and fine tune it

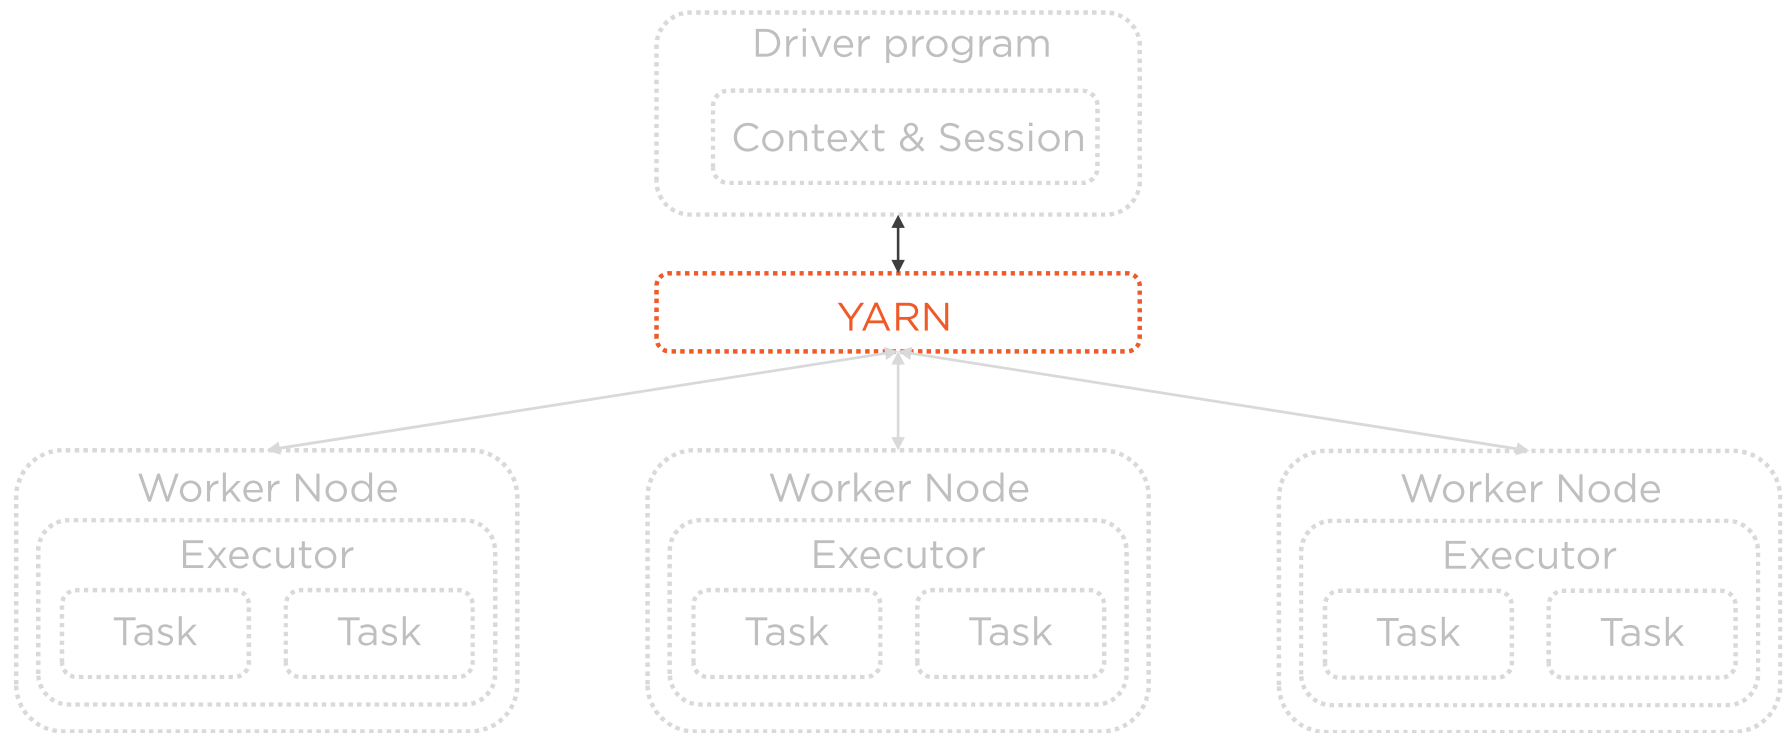
- Your requirements & resources
- Workload type




Spark's Architecture



Spark on YARN



YARN



Yet Another Resource Negotiator

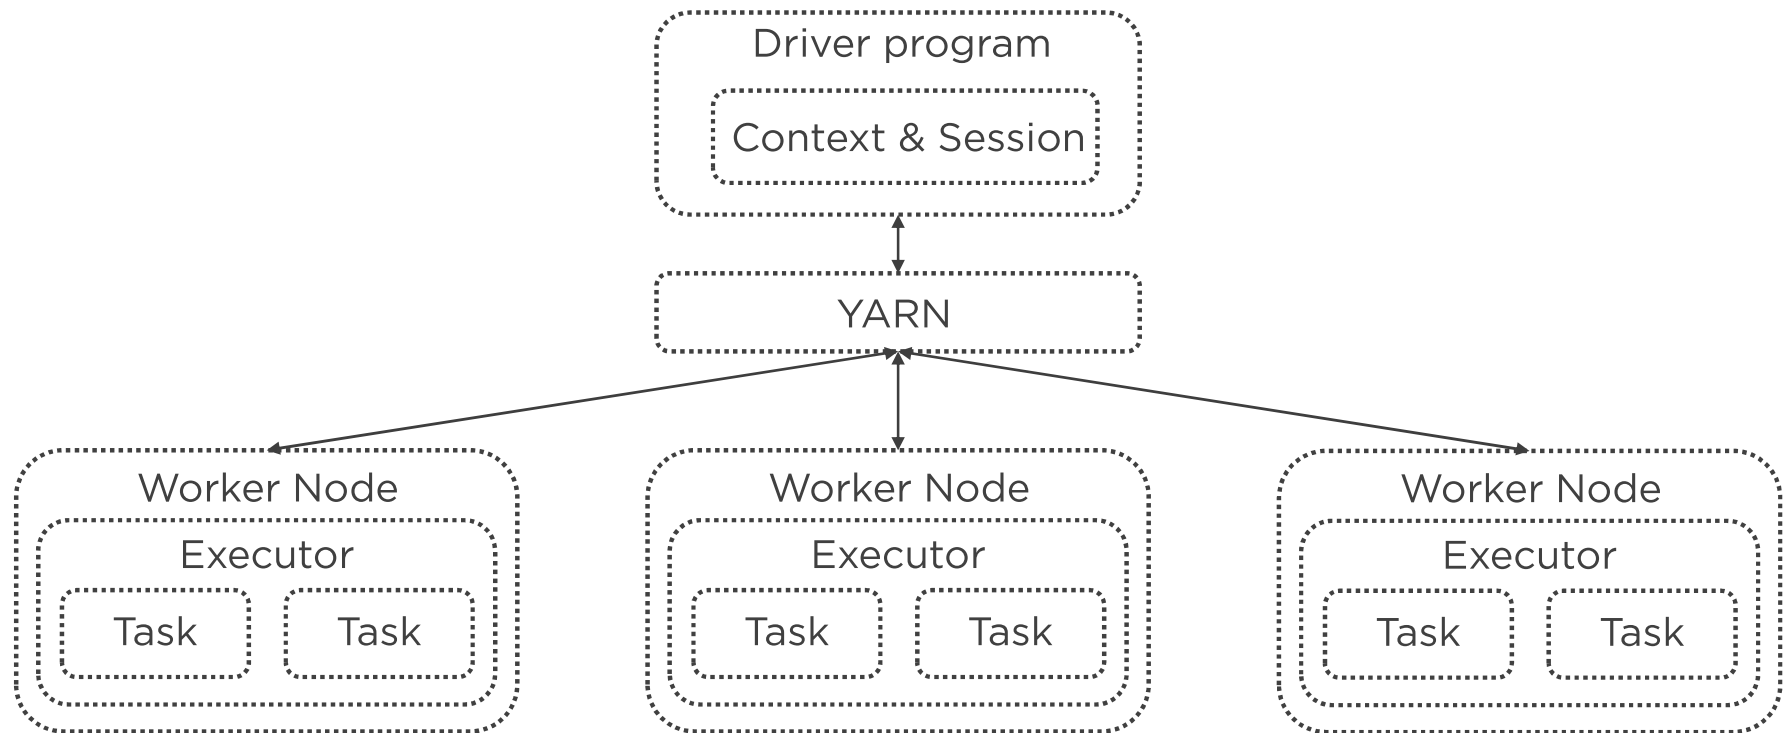
Cluster management technology

In charge of

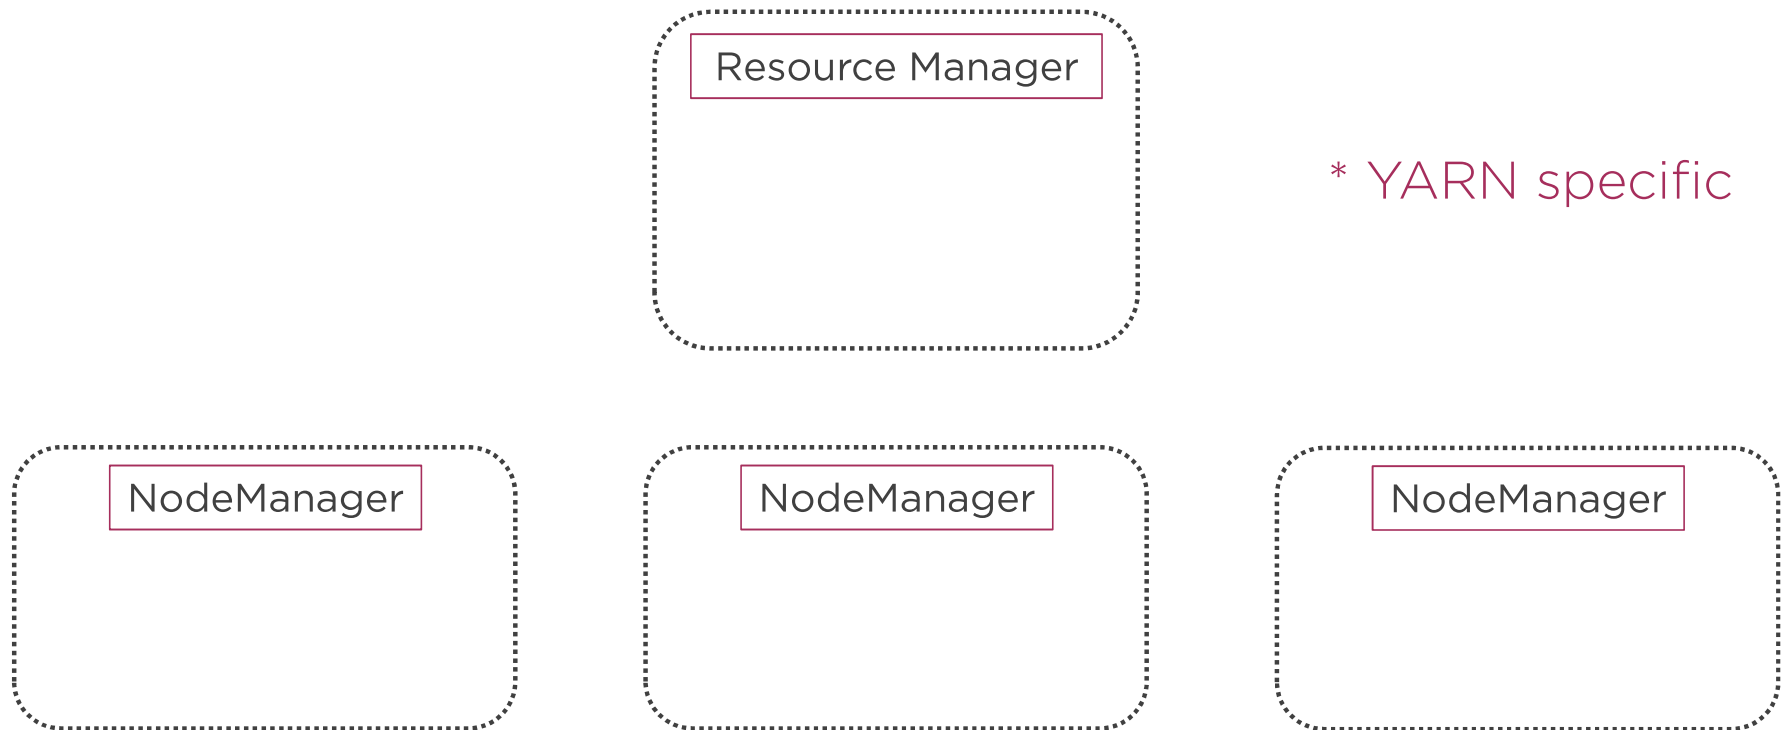
- Resource management
- Job scheduling/monitoring




Spark on YARN



Running Spark Jobs



Spark Cluster Mode



Resource Manager

NodeManager




NodeManager

Driver
Program

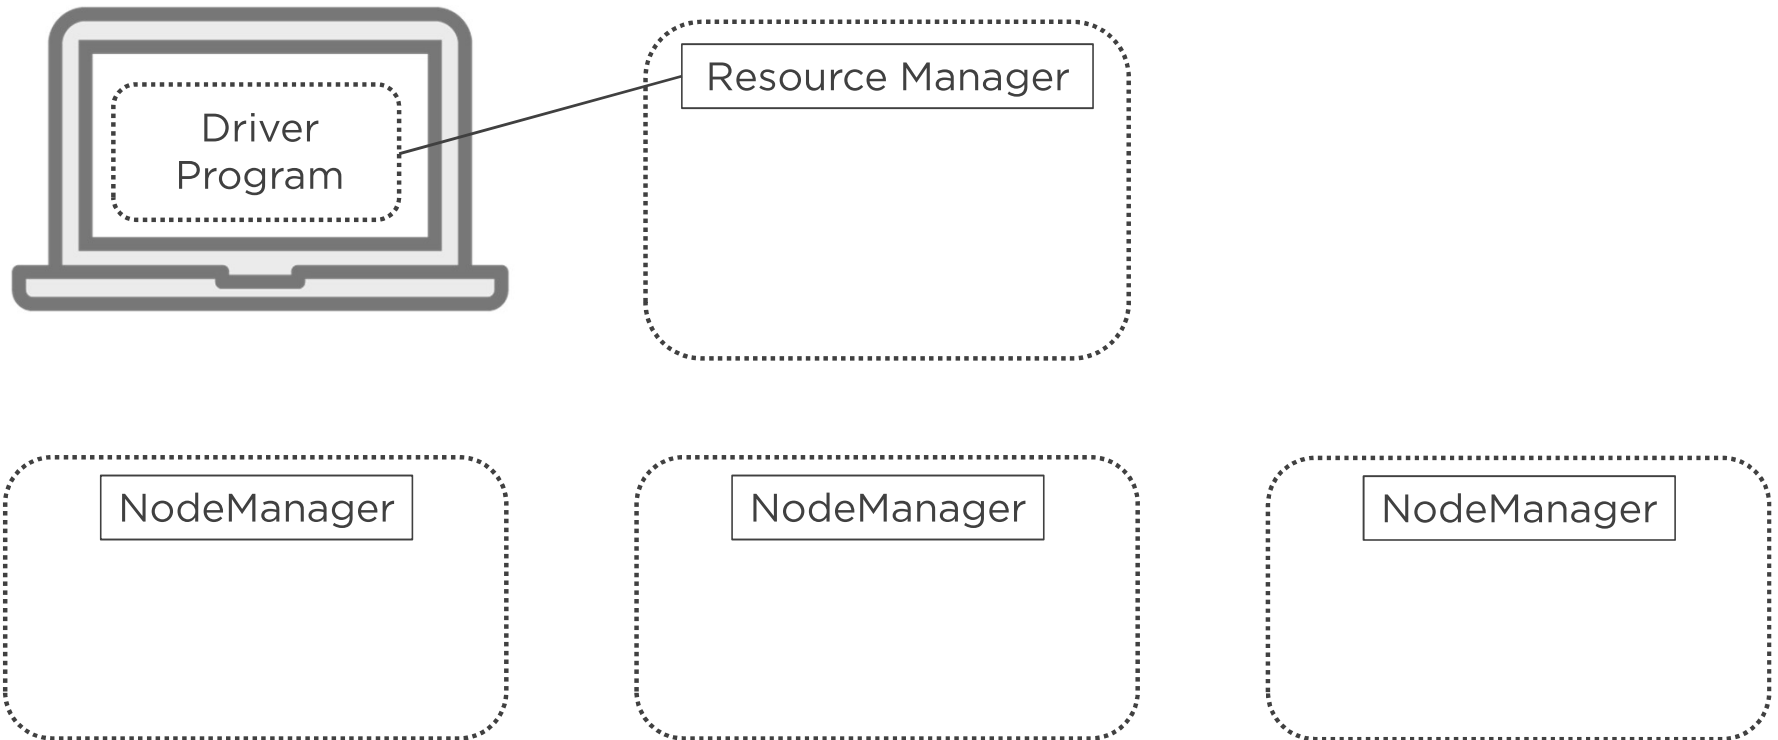
NodeManager




Spark Client Mode



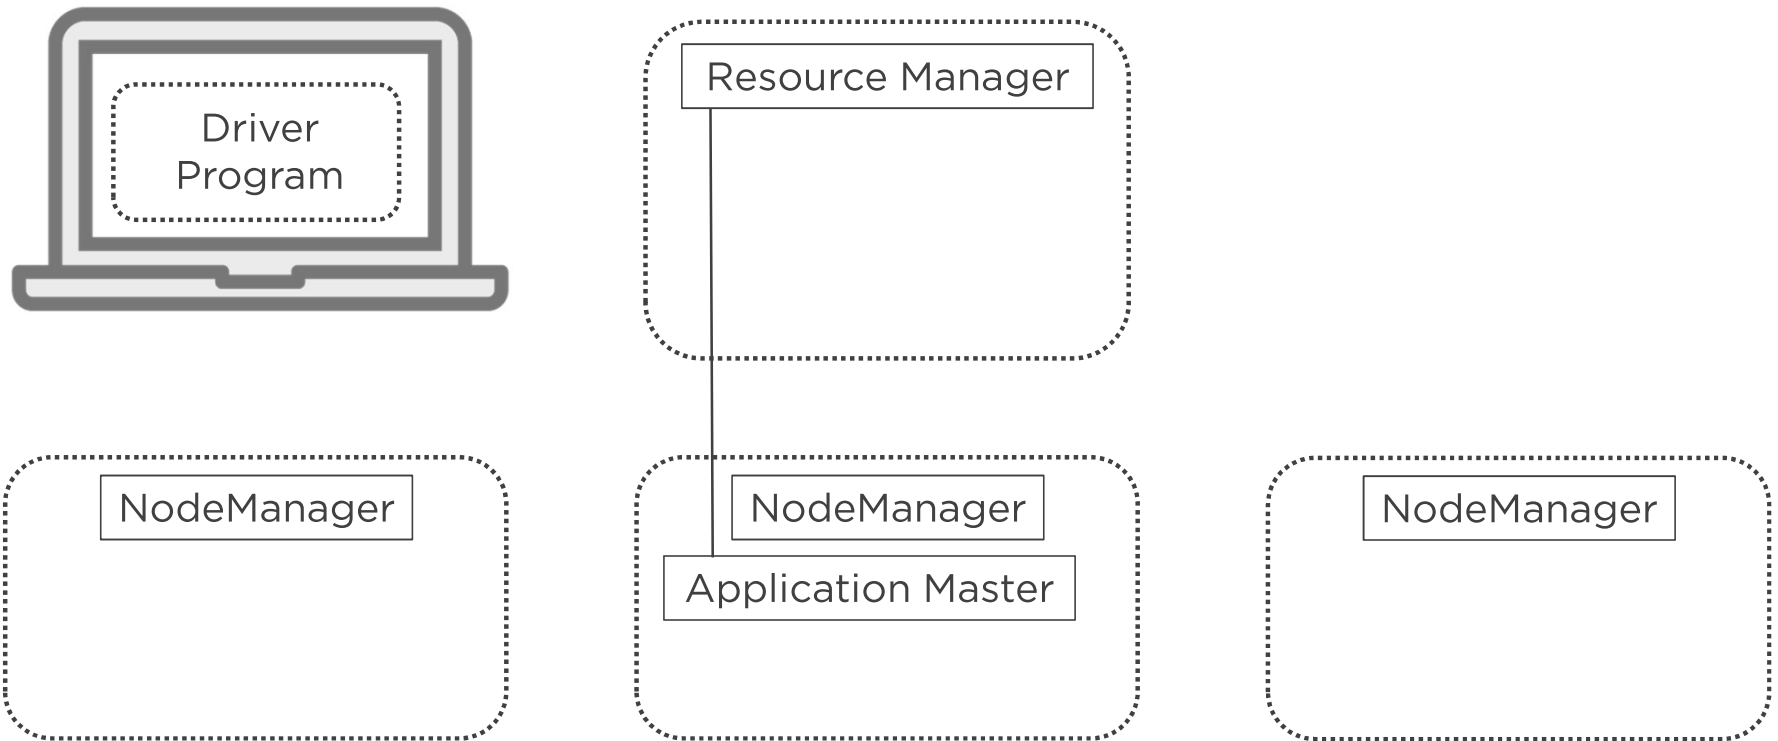
Spark Client Mode



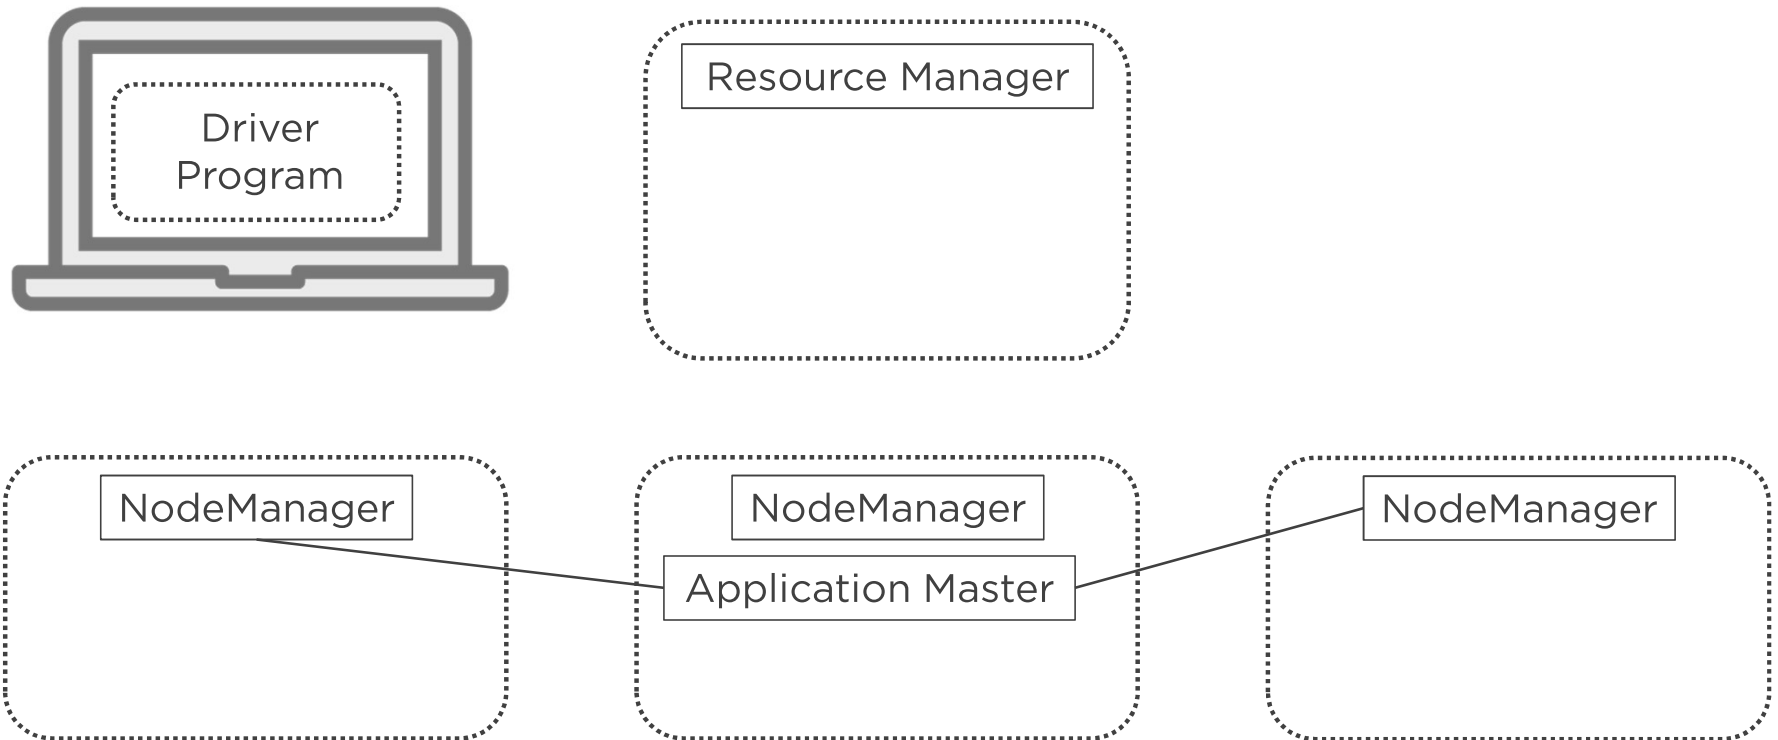
Spark Client Mode






Spark Client Mode



Spark Client Mode



Spark Client Mode



Home

Try Cloudera Enterprise Data Hub Edition for 60 Days

✓ Spark Pluralsight (CDH 5.12.1, Parcels) ▾

✓ Hosts

✓ HDFS ▾

✓ Hive ▾

✓ Hue ▾

✓ Oozie ▾

✓ Spark 2 ▾

✓ YARN (MR2...

Cloudera Management S

✓ Cloudera M...

Spark 2 Actions

Start

Restart

Stop





Instances

Configuration

Add Role Instances

Charts

30m 1h 2h 6h 12h 1d 7d 30d 



Visualizing Your Spark App: Web UI & History Server



Spark Jobs (?)

User: hdfs

Total Uptime: 41 min

Scheduling Mode: FIFO

Spark Web UI



History Server

Event log directory: hdfs://mgt01.cloudera:8020/user/

Last updated: 1/4/2018, 9:40:58 PM

Show entries

Spark History Server



Spark Web UI

Interface of a running Spark application

Launched by the SparkContext

Located in port 4040

- Configurable
- If port busy, will move to 4041 and on...



Spark Web UI

Available

- Timeline
- Resources
- DAG
- Storage
- Environment
- Executors
- SQL



Spark History Server

Web UI

History

- Completed apps
- Failed and killed apps
- Running apps

Port 18089



Spark History Server



Because Spark Web UI is removed after application exits



Log File

File that records events or messages



```
17/11/13 10:15:33 DEBUG ipc.ProtobufRpcEngine: Call: setTimes took 2ms
17/11/13 10:15:33 TRACE ipc.ProtobufRpcEngine: 35: Response <- mgt01.cloudera/10.0.2.100:8020: setTimes {}
17/11/13 10:15:33 INFO cluster.YarnClientSchedulerBackend: Interrupting monitor thread
17/11/13 10:15:33 TRACE client.TransportClient: Sending RPC to /10.0.2.100:52130
17/11/13 10:15:33 TRACE client.TransportClient: Sending request 7575887306779588512 to /10.0.2.100:52130 took 1 ms
17/11/13 10:15:33 TRACE protocol.MessageDecoder: Received message RpcResponse: RpcResponse {requestId=7575887306779588512, body=NettyManagedBuffer{buf=PooledUnsafeDirectByteBuf(ridx: 21, widx: 68, cap: 2048)}}
17/11/13 10:15:33 INFO cluster.YarnClientSchedulerBackend: Shutting down all executors
17/11/13 10:15:33 INFO cluster.YarnSchedulerBackend$YarnDriverEndpoint: Asking each executor to shut down
17/11/13 10:15:33 INFO cluster.SchedulerExtensionServices: Stopping SchedulerExtensionServices (serviceOption=None, services=List(), started=false)
17/11/13 10:15:33 DEBUG service.AbstractService: Service: org.apache.hadoop.yarn.client.api.impl.YarnClientImpl entered state STOPPED
17/11/13 10:15:33 DEBUG ipc.Client: stopping client from cache: org.apache.hadoop.ipc.Client@21461fc
17/11/13 10:15:33 INFO cluster.YarnClientSchedulerBackend: Stopped
17/11/13 10:15:33 INFO spark.MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
17/11/13 10:15:33 INFO memory.MemoryStore: MemoryStore cleared
17/11/13 10:15:33 INFO storage.BlockManager: BlockManager stopped
```

Logs

◀ 30 minutes preceding Nov 13, 10:44 AM EST ▶▶▶

Search

Additional Settings

Select Sources

Hosts

Minimum Log Level

INFO ▾

30m

1h

2h

6h

12h

1d

7d

30d


Previous

Next

Machines Searched: 3, Errors: 1, Search Time: 200 ms, [More Statistics](#)

Host	Log Level	Time	Source	Message
dn04.cloudera	INFO	November 13, 2017 10:19 AM	RMApManager\$ApplicationSummary	appId=application_1510324231586_0003, name=inExecutorEntry.py, user=hdfs, queue=root.users.hdfs, state=FINISHED, trackingUrl=http://dn04.cloudera:8088/proxy/application_1510324231586_0003/, appMasterHost=10.0.2.101, startTime=1510586320259, finishTime=1510586396900, finalStatus=SUCCEEDED, memorySeconds=169459, vcoreSeconds=135, preemptedAMContainers=0, preemptedNonAMContainers=0, preemptedResources=<memory:0\, vCores:0> View Log File
dn01.cloudera	INFO	November 13, 2017 10:37 AM	24235 CP Server Thread-6_cplogging	10.0.2.100 - - [13/Nov/2017:10:37:20] "GET /search_logs?start_time=1510585635462&end_time=1510587435462&log_level=INFO&query=inExecutorEntry&role_result_limit=101&log_paths=%2Fvar%2Flog%2Fhadoop-hdfs%2Fhadoop-cmf-hdfs-DATANODE-dn01.cloudera.log.out%2C%2Fvar%2Flog%2Fhive%2Fhadoop-cmf-hive-HIVEMETASTORE-dn01.cloudera.log.out%2C%2Fvar%2Flog%2Fhive%2Fhadoop-cmf-hive-HIVESERVER2-dn01.cloudera.log.out%2C%2Fvar%2Flog%2Fhadoop-yarn%2Fhadoop-cmf-yarn-NODEMANAGER-dn01.cloudera.log.out&log_types=LOG4J%2CLOG4J%2CLOG4J%2CLOG4J&role_ids=32%2C86%2C87%2C41&add_agent_logs=true&search_timeout=59979 HTTP/1.1" 200 305 "" "NING/1.0"

How Can I Find the Log Files?



Using the terminal, in Spark Log Dir

- Local
- HDFS

Spark UI

Spark History Server

Cloudera Manager

HUE

YARN





```
sc.setLogLevel("ALL")
```

Logging Level

Specify log4j's logging levels

Using **setLogLevel()**



Log Levels

OFF

No logging

FATAL

Only very severe errors, those that might cause application to abort

ERROR

Log errors, but that still may allow application to keep executing



Log Levels

WARN

Logs potentially harmful scenarios

INFO

Informational messages, that highlight progress of application

DEBUG

Used mostly by developers, mainly when debugging the application



Log Levels

TRACE

Even finer grained than DEBUG

ALL

Turns on all logging



Navigating the Documentation



Takeaway



Spark's architecture

- Driver, cluster manager, executor, tasks
- Storage
- Deployment modes
- YARN



Takeaway



Visualizing your Spark application

- Web UI
- History Server

Logging

Navigating the documentation



Takeaway



RDD API

