# Learning the Core of Spark: RDDs



Xavier Morera
HELPING DEVELOPERS UNDERSTAND SEARCH & BIG DATA
@xmorera www.xaviermorera.com



#### Learning the Core of Spark: RDDs



#### **Resilient Distributed Dataset**

How can I use DataFrames in 2.0

What is an RDD and Schema RDD

How do I group by a field

Can I use Hive from HUE

Datasets (DataFrames) & Spark SQL



## Why Start with RDDs?

Knowing RDDs gives you a better foundation for writing Spark Code

Sometimes RDDs are the right tool for the job

Existing Spark 1.x code



#### SparkContext Driver program Context Job Cluster Manager Worker Node Worker Node Worker Node Executor Executor Executor Task Task Task Task Task

You will need it to work with RDDs!



## SparkContext

#### **Spark Application**

- Created for you in the shell
- You need to create in an application

Only one SparkContext



# SparkContext



Spark Application



Configuration & Environment



Services



**Create RDDs** 



sc
spark.sparkContext
sc.sparkUser
sc.getConf.getAll

#### SparkContext

SparkContext available as sc

Many methods and properties available

Get configuration and environment information



```
sc.stop
sc
val test_rdd = sc.emptyRDD
test_rdd.collect()
import org.apache.spark.SparkContext
val sc = SparkContext.getOrCreate()
test_rdd = sc.emptyRDD
test_rdd.collect()
```

#### SparkContext

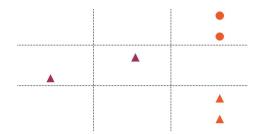
Stop and create a new SparkContext

- getOrCreate()

Used to create RDDs



#### RDD



How can I use <u>DataFrames</u> in 2.0

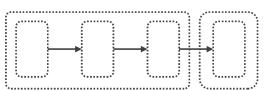
What is an RDD and Schema RDD

How do I group by a field

Can I use Hive from HUE

How can I use DataFrames in 2.0
What is an RDD and Schema RDD
How do I group by a field
Can I use Hive from HUE

Resilient



Distributed

**Dataset** 



## Five Properties of an RDD

Partitions

Partitioner (key/value RDDs)

Dependencies

Preferred locations for compute

Function to compute partition



```
* Internally, each RDD is characterized by five main properties:
62
63
64
     * - A list of partitions
     * - A function for computing each split
65
     * - A list of dependencies on other RDDs
66
     * - Optionally, a Partitioner for key-value RDDs (e.g. to say that the RDD is hash-partitioned)
67
     * - Optionally, a list of preferred locations to compute each split on (e.g. block locations for
68
69
          an HDFS file)
70
     * All of the scheduling and execution in Spark is done based on these methods, allowing each RDD
71
     * to implement its own way of computing itself. Indeed, users can implement custom RDDs (e.g. for
72
73
     * reading data from a new storage system) by overriding these functions. Please refer to the
74
     * <a href="http://people.csail.mit.edu/matei/papers/2012/nsdi_spark.pdf">Spark paper</a>
     * for more details on RDD internals.
75
76
     */
77
    abstract class RDD[T: ClassTag](
        @transient private var _sc: SparkContext,
78
79
        @transient private var deps: Seg[Dependency[_]]
      ) extends Serializable with Logging {
80
81
      if (classOf[RDD[_]].isAssignableFrom(elementClassTag.runtimeClass)) {
82
83
        // This is a warning instead of an exception in order to avoid breaking user programs that
        // might have defined nested RDDs without running jobs with them.
84
        logWarning("Spark does not support nested RDDs (see SPARK-5063)")
      }
```

PairRDD

RDDs of key/value pairs

**Contain tuples** 

Useful for grouping or aggregating

Can use RDD transformations

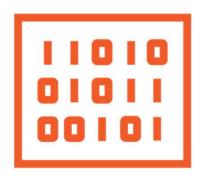
PairRDD specific transformations available



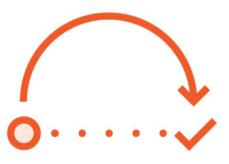
# Creating RDDS



**Parallelize** 



**External Data** 



From Another RDD

\* or Dataset

\* or DataFrame

(upcoming module)



```
Array(1,2,3,4,5)
val list_one_to_five = sc.parallelize(Array(1,2,3,4,5))
list_one_to_five
list_one_to_five.collect()
```

#### Creating RDDs with Parallelize

Create RDDs with data in memory: collection or iterable

Parallelize a collection

Bring data back to driver with collect()



```
list_one_to_five.getNumPartitions
list_one_to_five.glom().collect()
val list_one_to_five = sc.parallelize(Array(1,2,3,4,5), 1)
list_one_to_five.getNumPartitions
list_one_to_five.glom().collect()
```

#### RDD Partitions with Parallelize

Call getNumPartitions() to check number of partitions

Use glom() to see our data within partitions

And you can control how many partitions you want



```
list_one_to_five.sum()
list_one_to_five.min()
list_one_to_five.max()
list_one_to_five.mean()
list_one_to_five.first()
```

#### An Operation on an RDD

What can I do on an RDD?

Call sum() to add up values

And first() to get first element



```
val list_dif_types = sc.parallelize(Seq(false, 1, "two",
    Map("three" -> 3), ("xavier", 4)))
list_dif_types.collect()
val tuple_rdd = sc.parallelize(Seq(("xavier" -> 1), ("irene" -> 2)))
tuple_rdd
tuple_rdd.setName("tuple_rdd")
tuple_rdd
```

## Different Types of Objects

RDDs can hold different types of objects

Strings, ints, dictionaries, and key/value pairs (tuple) are common

RDD is given an internal name, but you can set it too



```
val empty_rdd = sc.parallelize(Array[String]())
empty_rdd.collect()
empty_rdd.isEmpty()
val other_empty = sc.emptyRDD
other_empty.isEmpty()
```

## EmptyRDD

You can create empty RDDs

With parallelize()

Or emptyRDD



```
val bigger_rdd = sc.parallelize(1 to 1000)
bigger_rdd.collect()
bigger_rdd.sum()
bigger_rdd.count()
bigger_rdd.getNumPartitions
```

#### Ranges

Create RDDs with a range (1 to 1000)

And use methods to get information from our RDD data

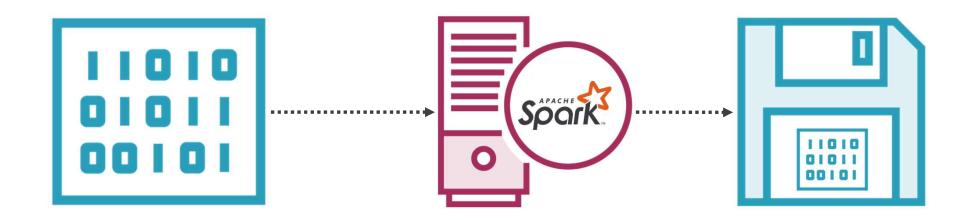
Count() all elements in my rdd

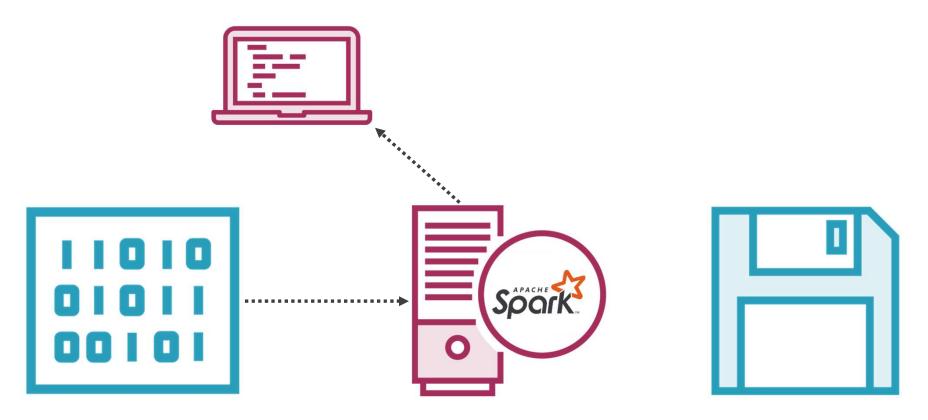


Parallelize

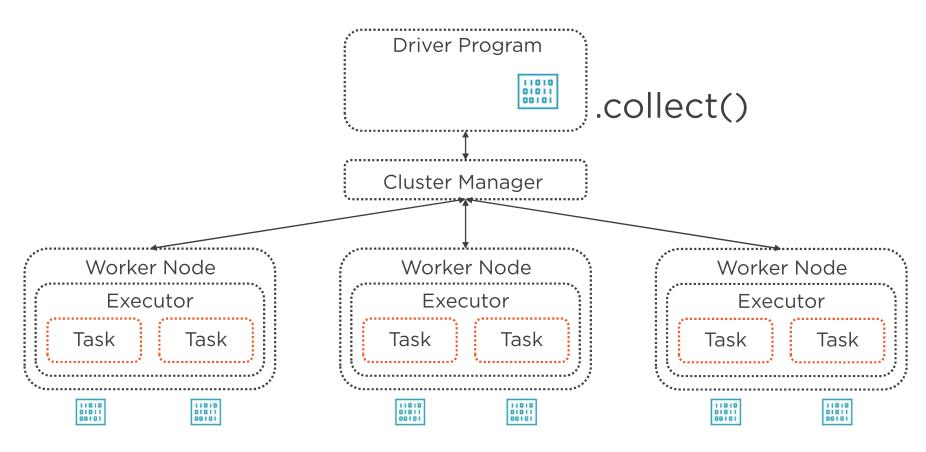












```
bigger_rdd.collect()
bigger_rdd.take(10)
bigger_rdd.first()
bigger_rdd.takeOrdered(10)(Ordering[Int].reverse)
val tuple_map = tuple_rdd.collectAsMap()
tuple_map("xavier")
```

Bring data back RDD to Driver

Most common action for testing is collect()

Also first(), take(), takeOrdered(), collectAsMap() among others



# ABeware!

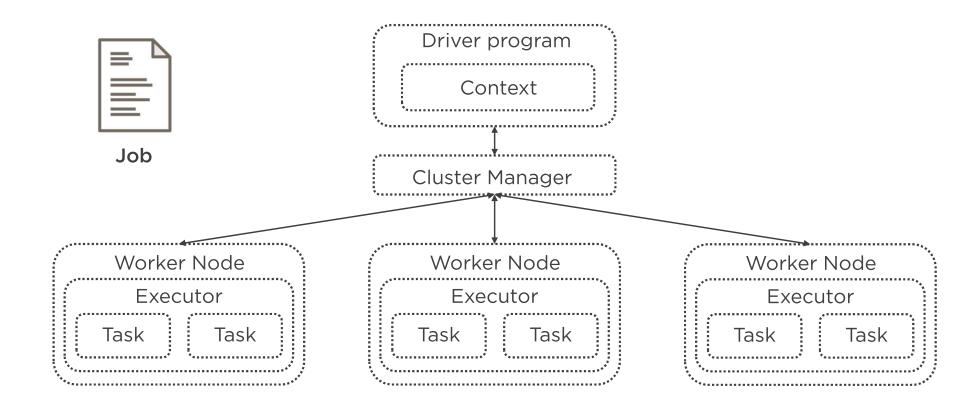
Potential out of memory exception on large datasets



```
for (elem <- bigger_rdd.take(10)) println(elem)
for ((k, v) <- tuple_map)
  println(k + " / " + v)</pre>
```

Iterating & Printing Data
Iterate over list of data returned using take() in for loop
Or other types, i.e. dict

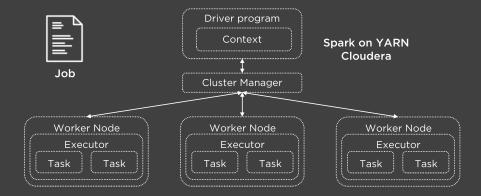




#### Driver

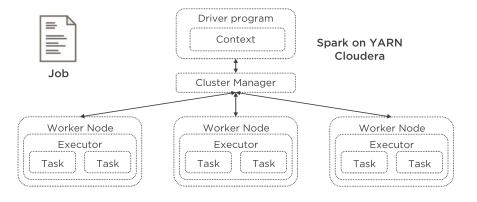
```
val test_write = sc.parallelize(
   Array("xavier", "troy"))
```

println(test\_write.collect()
 .mkString(" "))



#### In executor

```
def in_executor(entry: String) = {
  println("** inExecutorEntry **")
  println(entry)
  "printed: " + entry}
```



spark2-submit inExecutor EntryLogLevelAll.scala



```
def log_search(url: String): String = {
   val page =
   scala.io.Source.fromURL("http://tiny.bigdatainc.org/" +
   url).mkString
   println(page)
   page
}
val queries = sc.parallelize(Array("ts1", "ts2"))
queries.foreach(log_search)
```

#### Foreach

Executes the function on each element in the RDD

Runs in the worker nodes

Useful for calling external resources like a database or an API



```
bigger_rdd.getNumPartitions()
play_part=bigger_rdd.repartition(10)
play_part.getNumPartitions()
```

(Re)Partitions, Coalesce, Saving Text & HUE bigger\_rdd was partitioned on parallelize()

Change number of partitions with repartition() and coalesce()



play\_part.repartition(14).getNumPartitions
play\_part.coalesce(15).getNumPartitions

#### Repartition & Coalesce

repartition() specifies new number of partitions, up or down coalesce() uses existing partitions, only decreases, no shuffling Data in RDD remains the same. How can we tell?



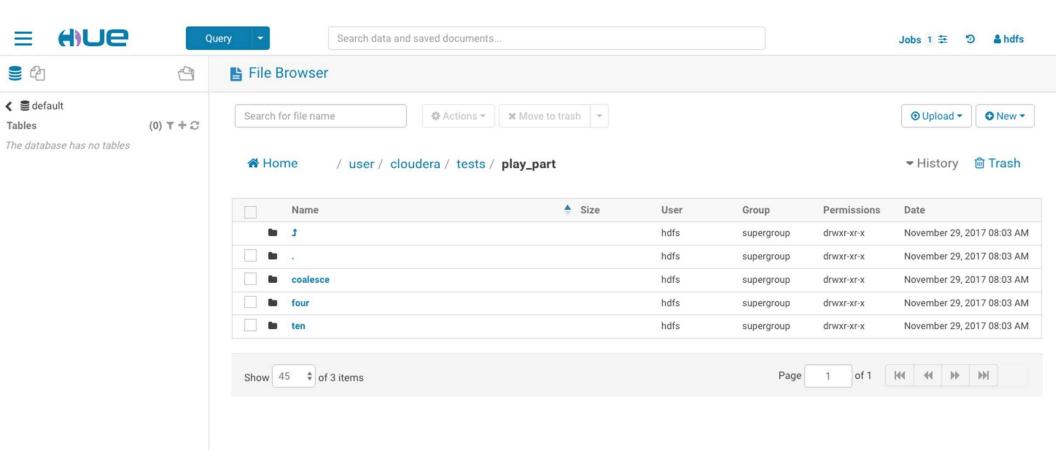
```
play_part.saveAsTextFile("/user/cloudera/tests/play_part/ten")
play_part.repartition(4).
    saveAsTextFile("/user/cloudera/tests/play_part/four")
play_part.coalesce(1).
    saveAsTextFile("/user/cloudera/tests/play_part/coalesce")
```

#### Confirm Data in RDD

Save using saveAsTextFile() and visualize with HUE

Just like Hadoop, creates file per partition







# 1. Warning!

Exception if folder already exists



```
play_part.coalesce(1).
   saveAsTextFile("/user/cloudera/tests/play_part/coalesce")
```

#### Folder Exists

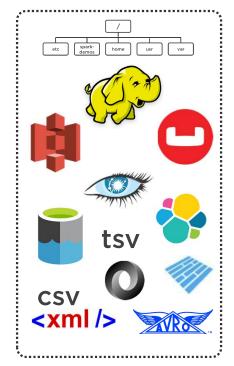
If folder exists, you get an exception

This applies for lower level API

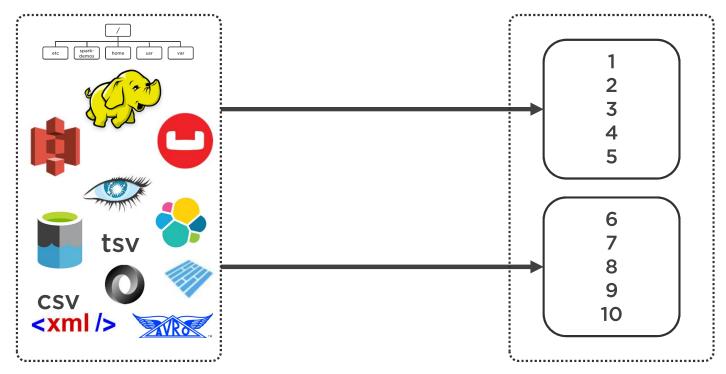
Works different for higher level API











```
sc.textFile("/user/cloudera/tests/play_part/ten").count()
sc.textFile("/user/cloudera/tests/play_part/four").count()
sc.textFile("/user/cloudera/tests/play_part/coalesce").count()
sc.textFile("/user/cloudera/tests/play_part/four/part-00000").count()
```

## Loading Text from HDFS

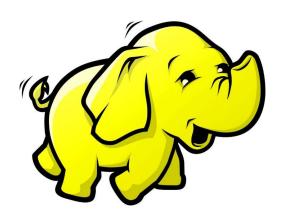
Use textFile(), using the SparkContext

Use hdfs:// with a different cluster manager

All files in the folder, or specific file



#### **HDFS**



**Hadoop Distributed File System** 

Primary storage system for Hadoop

**Data replication** 

**Blocks** 

Optimized for parallel data processing



```
val local_play =
   sc.textFile("file:///stackexchange/play_part/")
local_play.count
local_play.take(10)
```

## Loading Local Data

Access local data using file:///

File needs to be in all nodes

Not required in standalone Spark



#### Convert Posts.xml to CSV



# Data preparation step



```
posts_all =
sc.textFile("/user/cloudera/stackexchange/posts_all_csv")
posts_all.count()
```

Loading StackExchange / StackOverflow Posts

Load from HDFS

How many do we have?



#### CSV



**Comma Separated Values** 

Plain text file

**Tabular data** 

**Supports different separators** 

- Tab



#### Convert Badges.xml to CSV



# Data preparation step



```
def split_the_line(x: String): Array[String] = x.split(",")
val badges_rdd_csv =
    sc.textFile("/user/cloudera/stackexchange/badges_csv")
badges_rdd_csv.take(1)
badges_rdd_csv.take(1)(0)
val badges_columns_rdd =
    badges_rdd_csv.map(split_the_line)
badges_columns_rdd.take(1)(0)(2)
```

# Loading CSV

Loaded like any other file, with textFile()

You get lines, and tell Spark what to do with the data

Advantage and drawback at the same time



```
val numbers_partitions =
   sc.wholeTextFiles("/user/cloudera/tests/play_part/four")
numbers_partitions.take(1)
numbers_partitions.take(1)(0)._1
```

## Loading Whole Text Files

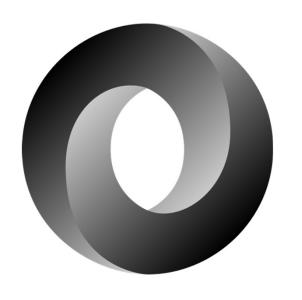
Loads each file into a row, uses wholeTextFiles()

In a PairRDD

Key is file name, value is the file



### **JSON**



**JavaScript Object Notation** 

Light weight data format

Easy to use

Collection of name value pairs and arrays

**Fat-free XML** 



#### Convert Tags.xml to JSON



# Data preparation step



```
val tags_json =
   sc.textFile("/user/cloudera/stackexchange/tags_json")
tags_json.first()
```

#### Load JSON

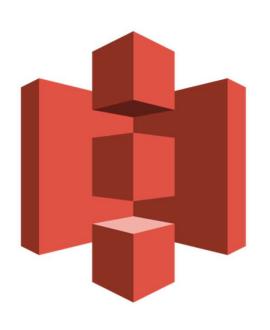
Just like CSV, it loads as text

Parse each record

Is there an easy way to parse JSON with RDDs?



#### S3



**Simple Storage Service** 

Amazon's cloud storage

**Buckets** 

**Rock solid SLA** 

**Unlimited scalability** 

At extremely reasonable cost

Works well with many technologies



Manually upload a JSON file into an S3 account
Get the full path in S3
Create an IAM user, enable programmatic access
Obtain access key id and secret access key
Set keys in environment
Load necessary packages

#### Prepare Data for Processing in S3



#### Data preparation step



```
# export AWS_ACCESS_KEY_ID="access-key"
# export AWS_SECRET_ACCESS_KEY="secret-key"

val tags_json_s3 = sc.textFile("s3a://pscs/part-00000")
tags_json_s3.take(10)
```

## Accessing Amazon S3 with Spark

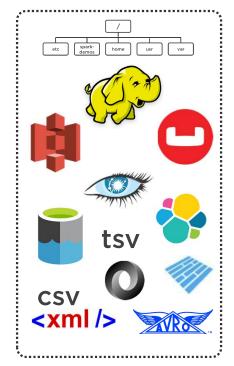
Transparent, still textFile()

Use s3a://

Be very careful with security







```
badges_columns_rdd.take(1)
badges_columns_rdd.take(1)(0)(2)
badges_columns_rdd.saveAsTextFile("/user/cloudera/
    stackexchange/badges_txt_array")
```

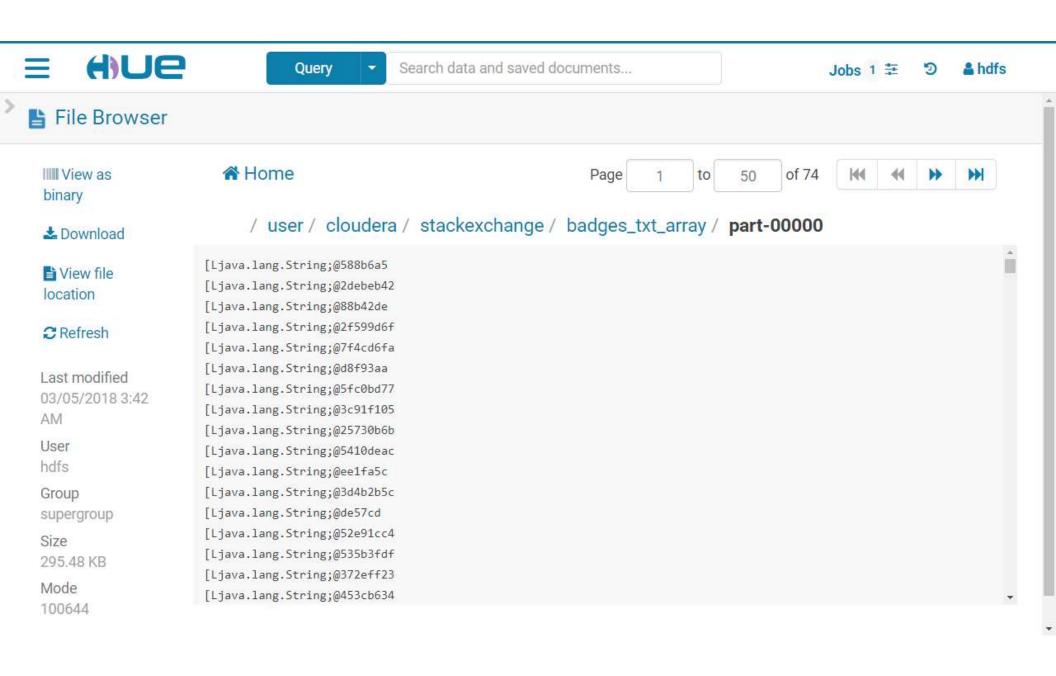
## Saving an RDD

We previously saved an RDD that had no transformations applied

- Partition files

Save an RDD with columns as text





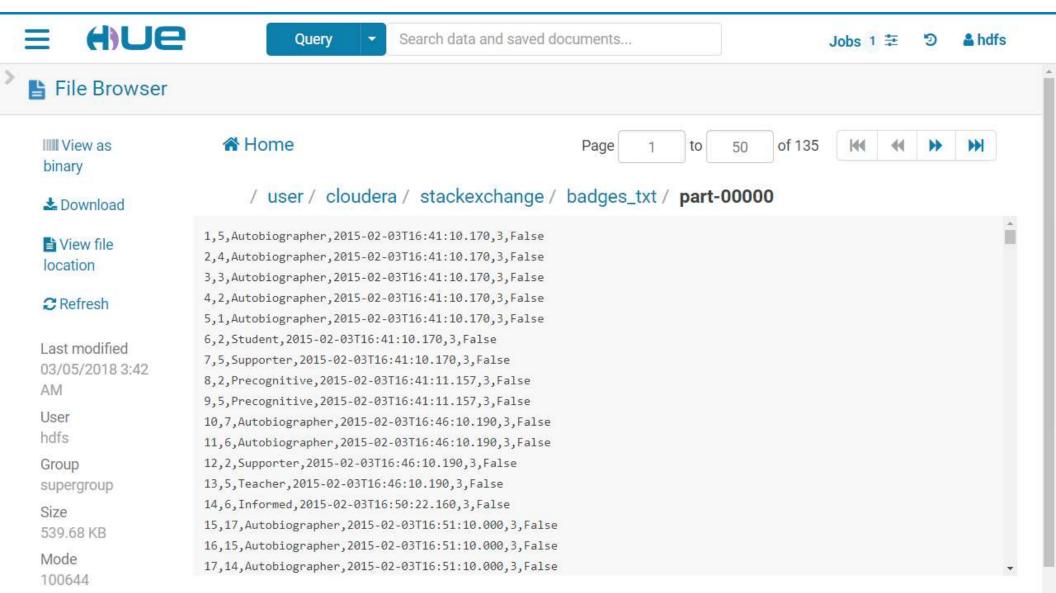
```
badges_columns_rdd.map(x => x.mkString(","))
    .saveAsTextFile("/user/cloudera/stackexchange/badges_txt")
```

### Saving an RDD

Save as text

Convert Array[String] to String





```
val badges_reloaded =
sc.textFile("/user/cloudera/stackexchange/badges_txt")
```

## Saving & Reloading an RDD

Reload our data

What happened when I save?



```
badges_reloaded.take(1)
badges_reloaded.take(1)(0)
badges_reloaded.take(1)(0)(0)
```

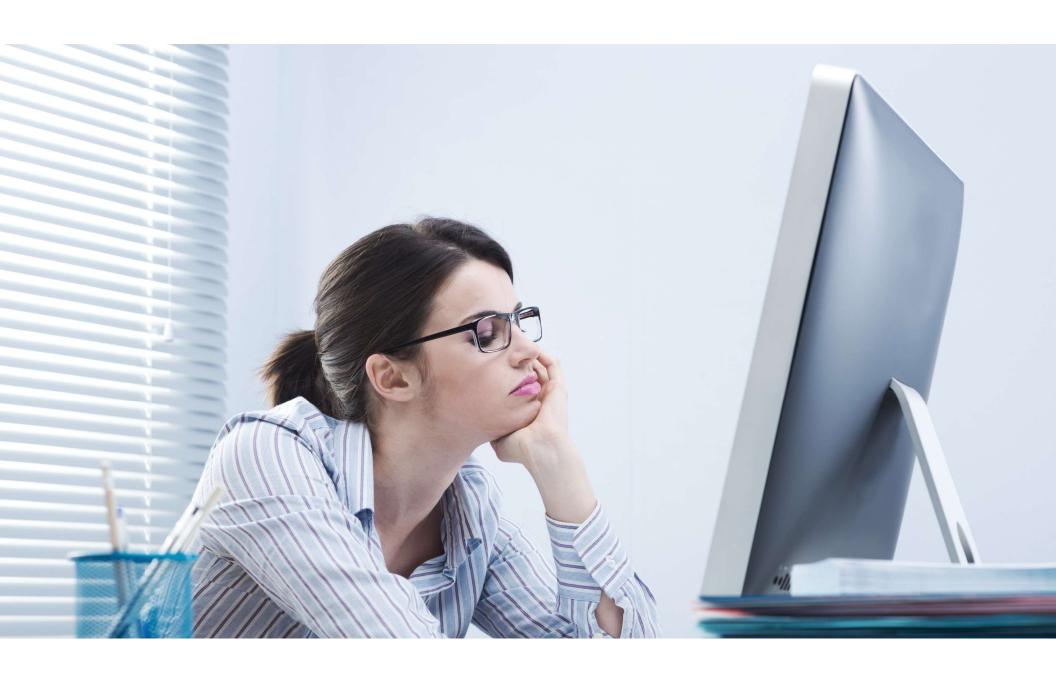
#### Text Is Good but Not Great

Compare original RDD vs. the one we just reloaded

We lost the format we applied

Text has no schema





# Object File



#### Save an RDD in a simple format

- ObjectFile

SequenceFile of serialized Java objects

Not the most efficient way

- Easy to use
- Still marked as experimental



```
badges_columns_rdd. //hit tab

badges_columns_rdd.saveAsObjectFile("/user/cloudera/
    stackexchange/badges_object")

val badges_object =
sc.objectFile[Array[String]]("/user/cloudera/stackexchange/
    badges_object")

badges_object.take(1)
```

## ObjectFile

Check available save methods

Save with saveAsObjectFile

And reload with objectFile() with type



# SequenceFile



Flat file

**Binary format** 

**Key/value pairs** 

**Used extensively in Hadoop** 

```
tuple_rdd.saveAsSequenceFile("/user/cloudera/stackexchange/
   tuple_sequence")

import org.apache.hadoop.io.Text

import org.apache.hadoop.io.LongWritable

sc.sequenceFile("/user/cloudera/stackexchange/tuple_sequence",
   classOf[Text], classOf[Longritable]).collect()
```

## SequenceFile

Save using saveAsSequenceFile()

Data needs to be in key/value format

Read with sequenceFile(), specify types



## Hadoop Formats



#### InputFormat

- How to read input files
- How they are split

#### OutputFormat

- How to write output files

#### Very flexible

```
badges_columns_rdd.map(prepareForNAH)
   .saveAsNewAPIHadoopFile("/user/cloudera/stackexchange/
    badgess_newapihadoop",
    classOf[Text], classOf[Text],
    classOf[SequenceFileOutputFormat[Text, Text]])
```

## Hadoop OutputFormat

Write using saveAsNewAPIHadopFile()

#### More parameters required than other methods

- i.e. file output format, key class, value class



## Hadoop OutputFormat

```
import org.apache.hadoop.mapreduce.lib.output._
import org.apache.hadoop.mapreduce.lib.input._
import org.apache.hadoop.io._
val prepareForNAH: Array[String] => (Text, Text) = x =>
(new Text(x(0)), new Text(x(2)))
badges_columns_rdd.map(prepareForNAH)
  .saveAsNewAPIHadoopFile("/user/cloudera/stackexchange/
    badgess_newapihadoop",
  classOf[Text],
  classOf[Text],
  classOf[SequenceFileOutputFormat[Text, Text]])
```



```
val badges_newapihadoop =
   sc.newAPIHadoopFile("/user/cloudera/stackexchange/
        badgess_newapihadoop",
   classOf[SequenceFileInputFormat[Text, Text]],
   classOf[Text],
   classOf[Text])

badges_newapihadoop.take(1)
```

## Hadoop InputFormat

Read with newAPIHadopFile()

Specify necessary parameters

Great deal of control



```
val xmlPosts =
  sc.newAPIHadoopFile("/user/cloudera/stackexchange/Posts.xml",
  classOf[TextInputFormat],
  classOf[LongWritable],
  classOf[Text])
```

Does newAPIHadoopFile Look Familiar?

We used to load the StackExchange dumps

Using spark-xml package





textFile
sequenceFile
objectFile
hadoopFile
hadoopDataset
newAPIHadoopFile
newAPIHdoopDataset



saveAsTextFile
saveAsSequenceFile
saveAsObjectFile
saveAsHadoopFile
saveAsHadoopDataset
saveAsNewAPIHadoopFile
saveAsNewAPIHdoopDataset





Full stop... nothing else to say...



```
val rdd_reuse = sc.parallelize(Array(1,2))
rdd_reuse.collect()
val rdd_reuse = sc.parallelize(Array(3,4))
rdd_reuse.collect()
```

### But... (On Immutable)

### **Immutable**

- Can't change the data

#### Reuse the variable



["1,5,Autobiographer,2015-02-03T00:00:00.000Z,3,False"]



["1", "5", "Autobiographer", "2015-02-03T00:00:00.000Z", "3", "False"]



```
badges_rdd_csv.take(1)
split_the_line
badges_columns_rdd.take(1)
```

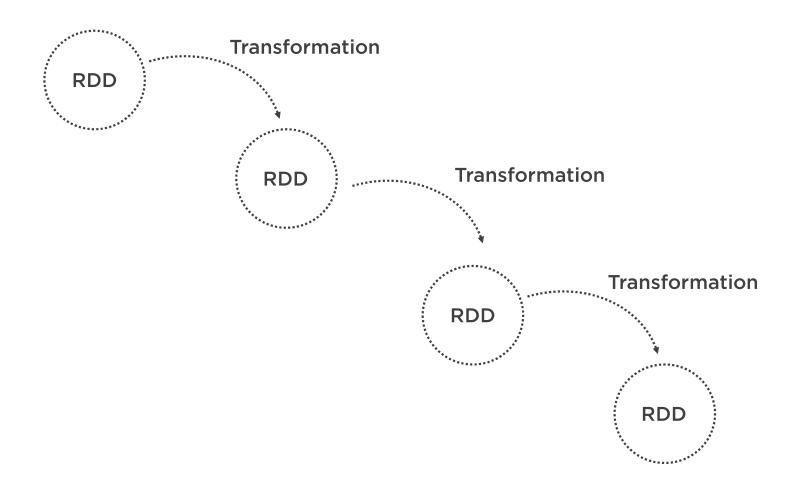
# Creating RDDs with Transformations

Loaded Badges CSV into an RDD

- Single element per row

Applied map(), with split\_the\_line() function





```
val badges_entry = badges_columns_rdd.map(x => x(2))
val badges_name = badges_entry.map(x => (x, 1))
val badges_reduced = badges_name.reduceByKey(_ + _)
val badges_count_badge = badges_reduced.
    map({ case (x,y) => (y,x) })
val badges_sorted = badges_count_badge.sortByKey(false).
    map({ case (x, y) => (y, x) })
```

# Creating Multiple RDDs with Transformations

#### One transformation

#### **Iterate**

- Applying multiple transformations
- One of the main objectives of Spark



```
$intp.definedTerms.
map(defTerms => s"${defTerms.toTermName}:
${$intp.typeOfTerm(defTerms.toTermName.toString)}").
filter(x => x.contains("()org.apache.spark.rdd.RDD")).
.foreach(println)
```

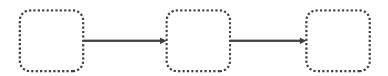
### Tip: Listing RDDs

Check for variables in memory with \$intp

Look for those that are an instance of RDD

- Use filter

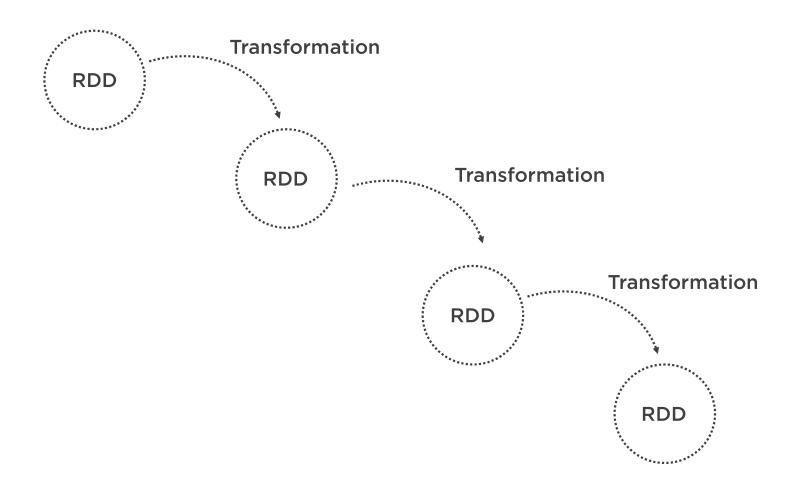




# RDD Lineage

Graph of transformation operations required to execute when an action is called

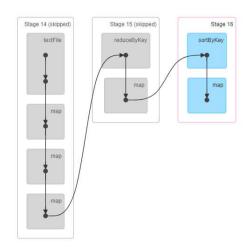
RDD operator graph or RDD dependency graph



badges\_sorted
badges\_sorted.toDebugString

A Little Bit More on Lineage Possible to see lineage

Use to Debug String()







Jobs

Stages Storage

Environment

Executors

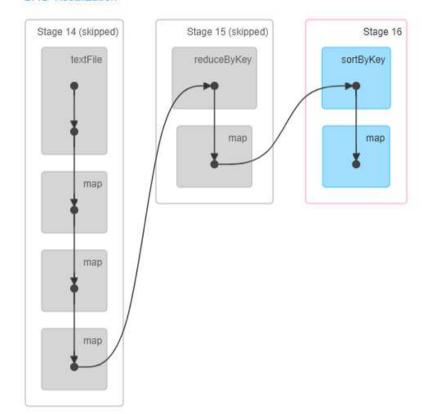
SQL

Spark shell application UI

#### **Details for Job 9**

Status: SUCCEEDED Completed Stages: 1 Skipped Stages: 2

- ▶ Event Timeline
- ▼ DAG Visualization



badges\_sorted.toDebugString
badges\_reduced.dependencies

## Dependencies

You \*could\* check RDD dependencies

At least to understand better the transformations





### Why learn RDDs?

- Better foundation
- Right tool for the job
- Existing code base



**SparkContext** 

One context per application

**RDD** main abstraction

Multiple types of RDDs





### **Create RDDs**

- Data in memory
- Files in storage
- Other RDDs

### Different data sources and file formats

- Text, CSV, Parquet, JSON, Sequence Files among others
- Transparent





### Return data back to the Driver

### Persist to disk

- Some formats store schema
- Others don't

### **RDD Lineage**

- Dependencies

