

Going Deeper into Spark Core



Xavier Morera

HELPING DEVELOPERS UNDERSTAND SEARCH & BIG DATA

@xmorera www.xaviermorera.com



Going Deeper into Spark Core



Why Lambdas?



Convenient and expressive

- Don't disregard named functions either

Quick to define, use and reuse

Testing and refactoring is easier

A.k.a Anonymous Functions



Anonymous Functions / Lambdas

Named
Functions
in Spark

```
def split_the_line(x: String):  
  Array[String] =  
    x.split(",")
```

```
badges.map(split_the_line)
```



Why Are Lambdas so Useful?

Anonymous
Functions
in Spark

~~badges.map(split_the_line)~~

badges.map(x => x.split(", "))



You will find yourself using
lambdas all the time with Spark

Believe me...



```
cd posts_simple_titles  
sbt package  
spark2-submit --class "PreparePostsSimpleTitlesApp"  
    target/scala-2.11/posts-simple-titles-project_2.11-1.0.jar
```

Extract Titles from Posts.xml



Data preparation step



```
val words_in_line = lines.map(x => x.split(" "))  
words_in_line.collect()
```

A Closer Look at Map, FlatMap, Filter, Sort, ...

map() is one of the most commonly used transformations

Followed by **flatMap()**, **filter()** and **sort()**

And later on aggregations




```
word_for_count = words.map(lambda x: (x,1))  
word_for_count.take(1)  
words.map(lambda x: x.lower())  
words.map(lambda x: x.upper())
```

Map

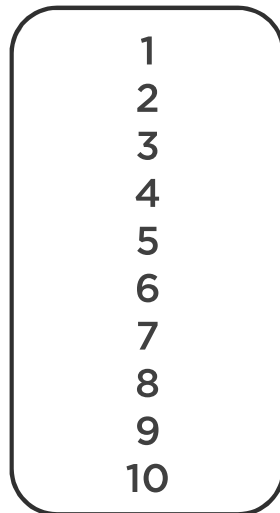
Apply function to each element

map(), mapPartitions(), mapValues(), mapPartitionsWithIndex() ...

RDD of length N transformed to RDD of length N



Parent RDD

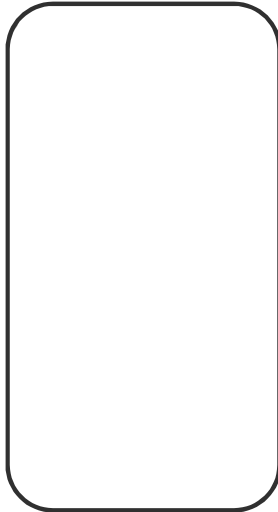


$x \Rightarrow x+1$

Child RDD

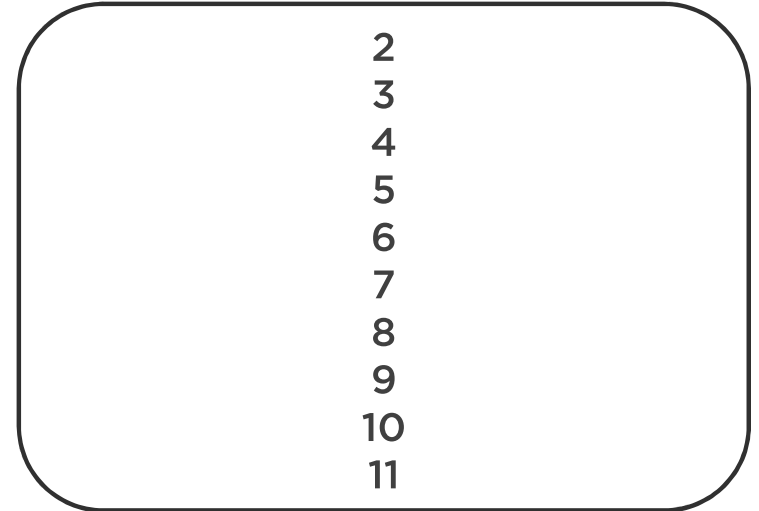


Parent RDD



$x \Rightarrow x+1$

Child RDD



```
val word_for_count = words.map(x => (x,1))  
word_for_count.take(1)  
word_for_count.take(5)
```

Map

Apply function to each element

map(), mapPartitions(), mapValues(), mapPartitionsWithIndex() ...

Each element in parent RDD mapped to one element in the child RDD



```
val words = lines.flatMap(line => line.split(" "))  
words.collect()
```

FlatMap

Apply function to each element and returns list of elements

Returns 0, 1 or more elements, "flattens" the results with map



Parent RDD

How can I use DataFrames in 2.0
What is an RDD and Schema RDD
How do I group by a field
Can I use Hive from HUE

Child RDD



Parent RDD



Child RDD

How, can, I, use, DataFrames,
in, 2.0, What, is, an, RDD, and,
Schema, RDD, How, do, I,
group, by, a, field, Can, I, use,
Hive, from, HUE



```
def starts_h(word: (String, Int)) =  
    word._1.toLowerCase.startsWith("h")  
  
word_for_count.filter(starts_h).collect()
```

Filter

Apply a function to each element of the RDD

If the function returns false, element is not included in new RDD

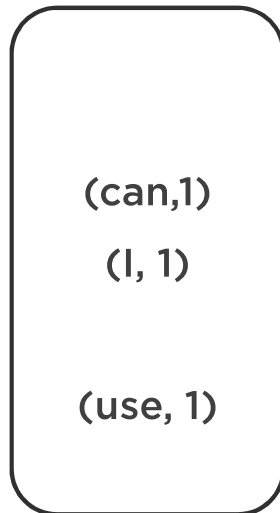


Filter

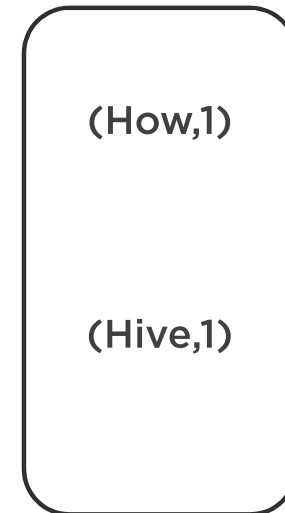


Filter

Parent RDD



Child RDD



```
val word_count = word_for_count.reduceByKey(_ + _)  
word_count.sortByKey().collect()  
word_count.sortByKey(false).collect()  
word_count.map({ case (x,y) => (y,x) }).sortByKey()  
    .map(x => x.swap).collect()  
word_count.sortBy({ case (x,y) => -y }).collect()
```

SortBy and SortByKey

Sort elements of an RDD

- By key on PairRDD with `sortByKey()`
- By a function using `sortBy()`



```
word_for_count.distinct().filter(starts_h).collect()
```

Many More Transformations

Plenty of transformations to go around

Some of them very powerful and/or very useful



Plenty of transformations to go around...



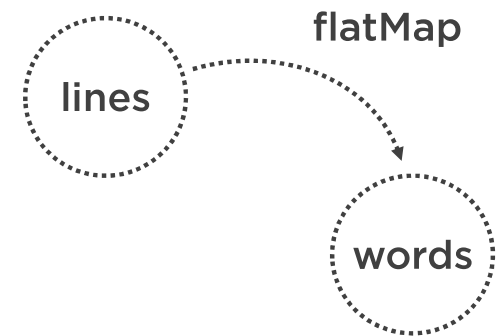
```
val lines = sc.textFile("/user/cloudera/se/simple_titles.txt")
val words = lines.flatMap(line => line.split(" "))
val word_for_count = words.map(x => (x,1))
val grouped_words = word_for_count.reduceByKey(_ + _)
grouped_words.collect()
```

Transformations

Start with method from **SparkContext** to load data

Transformations perform a computation

And create new RDDs



def **keyBy**[K](f: (T) ⇒ K): [RDD](#)[(K, T)]
Creates tuples of the elements in this RDD by applying f.

▶ def **localCheckpoint**(): [RDD](#).this.type
Mark this RDD for local checkpointing using Spark's existing caching layer.

def **map**[U](f: (T) ⇒ U)(*implicit* arg0: ClassTag[U]): [RDD](#)[U]
Return a new RDD by applying a function to all elements of this RDD.

▶ def **mapPartitions**[U](f: (Iterator[T]) ⇒ Iterator[U],
preservesPartitioning: Boolean = false)(*implicit* arg0:
ClassTag[U]): [RDD](#)[U]
Return a new RDD by applying a function to each partition of this RDD.

▶ def **mapPartitionsWithIndex**[U](f: (Int, Iterator[T]) ⇒ Iterator[U],
preservesPartitioning: Boolean = false)(*implicit* arg0:
ClassTag[U]): [RDD](#)[U]
Return a new RDD by applying a function to each partition of this RDD, while
tracking the index of the original partition.

▶ def **max**()(*implicit* ord: Ordering[T]): T
Returns the max of this RDD as defined by the implicit Ordering[T].

▶ def **min**()(*implicit* ord: Ordering[T]): T

Transformations

groupBy cartesian
flatMap intersection
map filter sample repartition union subtract
keyBy sortBy
coalesce zipWithIndex
zip mapPartitions
distinct



Transformations

PairRDDs

combineByKey sampleByKey

reduceByKey

join

leftOuterJoin

fullOuterJoin

sortByKey

cogroup

subtractByKey

groupByKey

rightOuterJoin

aggregateByKey

flatMapValues

foldByKey

reduceByKeyLocally

partitionBy



```
val lines = sc.textFile("/user/cloudera/se/simple_titles.txt")
val words = lines.flatMap(line => line.split(" "))
val word_for_count = words.map(x => (x,1))
val grouped_words = word_for_count.reduceByKey(_ + _)
grouped_words.collect()
```

Previously on Transformations

Transformations are what "changes" your data

Remember: Spark is lazy

No computation done when you specify transformation



```
val lines = sc.textFile("/user/cloudera/se/simple_titles.txt")
val words = lines.flatMap(line => line.split(" "))
val word_for_count = words.map(x => (x,1))
val grouped_words = word_for_count.reduceByKey(_ + _)
grouped_words.collect()
grouped_words.saveAsTextFile("/user/cloudera/stackexchange/words")
```

Actions

Action triggers computation

i.e. can return data to the driver or save an RDD to storage

Operations that produce non RDD values



RDD - org.apache.spark.rdd

← → ↺ spark.apache.org/docs/2.2.0/api/scala/index.html#org.apache.spark.rdd.RDD ☆

Q

#ABCDEFGHIJKLMNOPQRSTUVWXYZ

— deprecated

display packages only

hide focus

org.apache.spark.ui.env

EnvironmentListener

hide focus

org.apache.spark.ui.exec

ExecutorsListener

hide focus

org.apache.spark.ui.jobs

JobProgressListener

hide focus

org.apache.spark.ui.storage

StorageListener

hide focus

org.apache.spark.util

AccumulatorV2

CollectionAccumulator

DoubleAccumulator

EnumUtil

LegacyAccumulatorWrapper

LongAccumulator

MutablePair

SizeEstimator

StatCounter

TaskCompletionListener

TaskFailureListener

hide focus

org.apache.spark.util.random

BernoulliCellSampler

BernoulliSampler

PoissonSampler

Pseudorandom

Return the Cartesian product of this RDD and another one, that is, the RDD of all pairs of elements (a, b) where a is in this and b is in other.

def checkpoint(): Unit
Mark this RDD for checkpointing.

def coalesce(numPartitions: Int, shuffle: Boolean = false, partitionCoalescer: Option[PartitionCoalescer] = Option.empty)(implicit ord: Ordering[T] = null): RDD[T]
Return a new RDD that is reduced into numPartitions partitions.

def collect[U](f: PartialFunction[T, U])(implicit arg0: ClassTag[U]): RDD[U]
Return an RDD that contains all matching values by applying f.

def collect(): Array[T]
Return an array that contains all of the elements in this RDD.

Note This method should only be used if the resulting array is expected to be small, as all the data is loaded into the driver's memory.

def context: SparkContext
The org.apache.spark.SparkContext that this RDD was created on.

def count(): Long
Return the number of elements in the RDD.

def countApprox(timeout: Long, confidence: Double = 0.95): PartialResult[BoundedDouble]
Approximate version of count() that returns a potentially incomplete result within a timeout, even if not all tasks have finished.

def countApproxDistinct(relativeSD: Double = 0.05): Long
Return approximate number of distinct elements in the RDD.

def countApproxDistinct(p: Int, sp: Int): Long
Return approximate number of distinct elements in the RDD.

def countByValue()(implicit ord: Ordering[T] = null): Map[T, Long]
Return the count of each unique value in this RDD as a local map of (value, count) pairs.

def countByValueApprox(timeout: Long, confidence: Double = 0.95)(implicit ord: Ordering[T] = null): PartialResult[Map[T, BoundedDouble]]

PairRDDFunctions - org

spark.apache.org/docs/2.2.0/api/scala/index.html#org.apache.spark.rdd.PairRDDFunctions

display packages only

org.apache.spark.ui.env
EnvironmentListener

org.apache.spark.ui.exec
ExecutorsListener

org.apache.spark.ui.jobs
JobProgressListener

org.apache.spark.ui.storage
StorageListener

org.apache.spark.util
AccumulatorV2
CollectionAccumulator
DoubleAccumulator
EnumUtil
LegacyAccumulatorWrapper
LongAccumulator
MutablePair
SizeEstimator
StatCounter
TaskCompletionListener
TaskFailureListener

org.apache.spark.util.random
BernoulliCellSampler
BernoulliSampler
PoissonSampler
Pseudorandom

def cogroup[W1, W2, W3](other1: RDD[(K, W1)], other2: RDD[(K, W2)], other3: RDD[(K, W3)]): RDD[(K, (Iterable[V], Iterable[W1], Iterable[W2], Iterable[W3]))]
For each key k in this or other1 or other2 or other3, return a resulting RDD that contains a tuple with the list of values for that key in this, other1, other2 and other3.

def cogroup[W1, W2](other1: RDD[(K, W1)], other2: RDD[(K, W2)], partitioner: Partitioner): RDD[(K, (Iterable[V], Iterable[W1], Iterable[W2]))]
For each key k in this or other1 or other2, return a resulting RDD that contains a tuple with the list of values for that key in this, other1 and other2.

def cogroup[W](other: RDD[(K, W)], partitioner: Partitioner): RDD[(K, (Iterable[V], Iterable[W]))]
For each key k in this or other, return a resulting RDD that contains a tuple with the list of values for that key in this as well as other.

def cogroup[W1, W2, W3](other1: RDD[(K, W1)], other2: RDD[(K, W2)], other3: RDD[(K, W3)], partitioner: Partitioner): RDD[(K, (Iterable[V], Iterable[W1], Iterable[W2], Iterable[W3]))]
For each key k in this or other1 or other2 or other3, return a resulting RDD that contains a tuple with the list of values for that key in this, other1, other2 and other3.

def collectAsMap(): Map[K, V]
Return the key-value pairs in this RDD to the master as a Map.

Warning: this doesn't return a multimap (so if you have multiple values to the same key, only one value per key is preserved in the map returned)

Note this method should only be used if the resulting data is expected to be small, as all the data is loaded into the driver's memory.

def combineByKey[C](createCombiner: (V) => C, mergeValue: (C, V) => C, mergeCombiners: (C, C) => C): RDD[(K, C)]
Simplified version of combineByKeyWithClassTag that hash-partitions the resulting RDD using the existing partitioner/parallelism level.

def combineByKey[C](createCombiner: (V) => C, mergeValue: (C, V) => C, mergeCombiners: (C, C) => C, numPartitions: Int): RDD[(K, C)]
Simplified version of combineByKeyWithClassTag that hash-partitions the output RDD.

def combineByKey[C](createCombiner: (V) => C, mergeValue: (C, V) => C, mergeCombiners: (C, C) => C, partitioner: Partitioner, mapSideCombine: Boolean = true, serializer: Serializer = null): RDD[(K, C)]
Generic function to combine the elements for each key using a custom set of aggregation functions

Actions

histogram
saveAsHadoopDataset
collectAsMap
saveAsNewAPIHadoopDataset
collect
count
max
top
variance
mean
aggregate
fold
forEachPartition
saveAsHadoopFile
sampleVariance
countApprox
sum
takeSample
first
take
saveAsSequenceFile
saveAsObjectFile
treeReduce
saveAsNewAPIHadoopFile
treeAggregate
countApproxDistinct
saveAsTextFile
takeOrdered
min
stdev



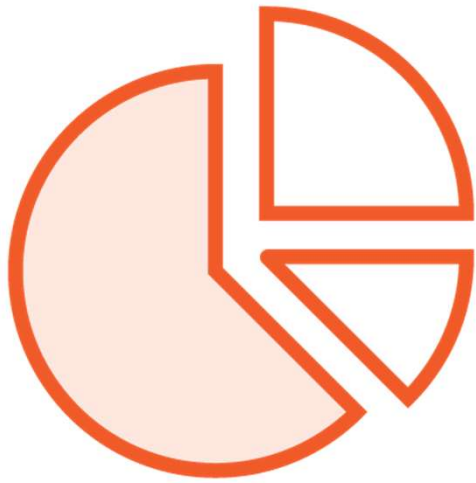
Actions
PairRDD

countApproxDistinctByKey

keys
values
countByKey
countByValueApprox
countByKeyApprox
sampleByKeyExact
countByValue



A Thing or Two on Partitions



Partition is just a 'bunch' of data

One of the foundations of parallelism

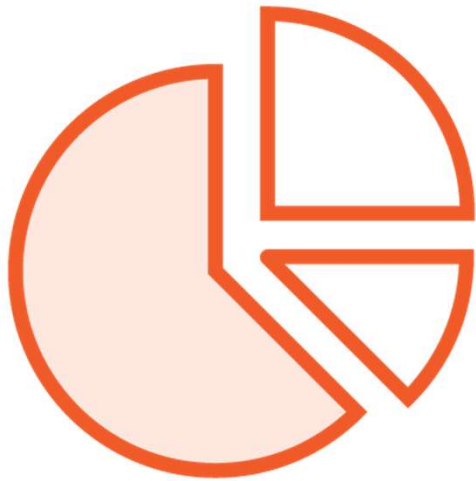
Faster to operate within partition

- Than shuffling data

Group data to minimize network traffic



How Does Spark Partition Data?



Data locality

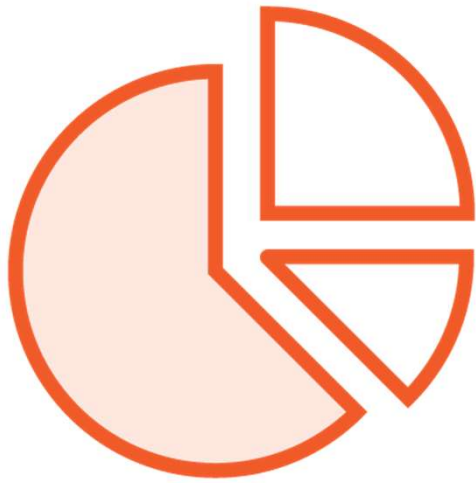
- Partition per HDFS block

Resources

Configuration or parameters



How Does Spark Partition Data?



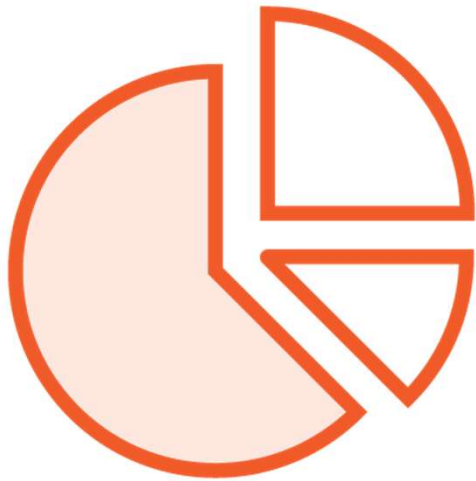
Partitioner

- Hash partitioner
- Range partitioner

Repartition



More or Less Partitions?



More partitions

- Less data per partition
- Smaller jobs
- More parallelism

Less partitions

- More data per partition
- Larger jobs



```
val badges_for_part = badges_columns_rdd.map(x => (x(2),  
x.mkString(", "))).repartition(50)  
  
badges_for_part.partitioner  
  
import org.apache.spark.HashPartitioner
```

PartitionBy

Returns an RDD partitioned using a specific **partitioner**

Useful to get keyed data into same partition

Not yet a group operation



```
val badges_by_badge = badges_for_part.  
    partitionBy(new HashPartitioner(50))  
  
badges_by_badge.partitioner  
  
badges_for_part.saveAsTextFile("/user/cloudera/  
    stackexchange/badges_no_partitioner")  
  
badges_by_badge.saveAsTextFile("/user/cloudera/  
    stackexchange/badges_yes_partitioner")
```

PartitionBy

Create a function to be used for partitioning

Pass function as parameter to **partitionBy()**

Save with and without partitioner, and review results



```
badges_by_badge.map({ case (x,y) => x }).glom().take(1)
```

Glom

There is an action to coalesce all rows in a partition into an array

Useful for operations on all items within a partition

Let's print our keys per partition



```
badges_by_badge.  
  mapPartitions(x => Array(x.size).iterator, true)  
  .collect()  
  
badges_for_part.  
  mapPartitions(x => Array(x.size).iterator, true)  
  .collect()
```

MapPartitions

Apply a function to each partition

Done at a single pass

Returns after entire partition is processed



```
posts_all.count()
```

Sampling Data

Selecting a representative part of the population

Faster, but you may lose accuracy

Also useful if you are resource constrained or very large dataset




```
val sample_posts = posts_all.sample(false, 0.1, 50)
sample_posts.count()
```

Sampling Data

Transformation to obtain a sample from your data with `sample()`

- withReplacement
- fraction
- seed



```
posts_all.count()
```

```
posts_all.countApprox(100, 0.95)
```

Approximate Counts

Obtain an approximate count with **countApprox()**

Note: Experimental



```
posts_all.takeSample(false, 15, 50)  
posts_all.takeSample(false, 15, 50).size
```

Take a Sample of Exact Size

Action available for exact count is called **takeSample()**



Set Operations



Questions

Answers



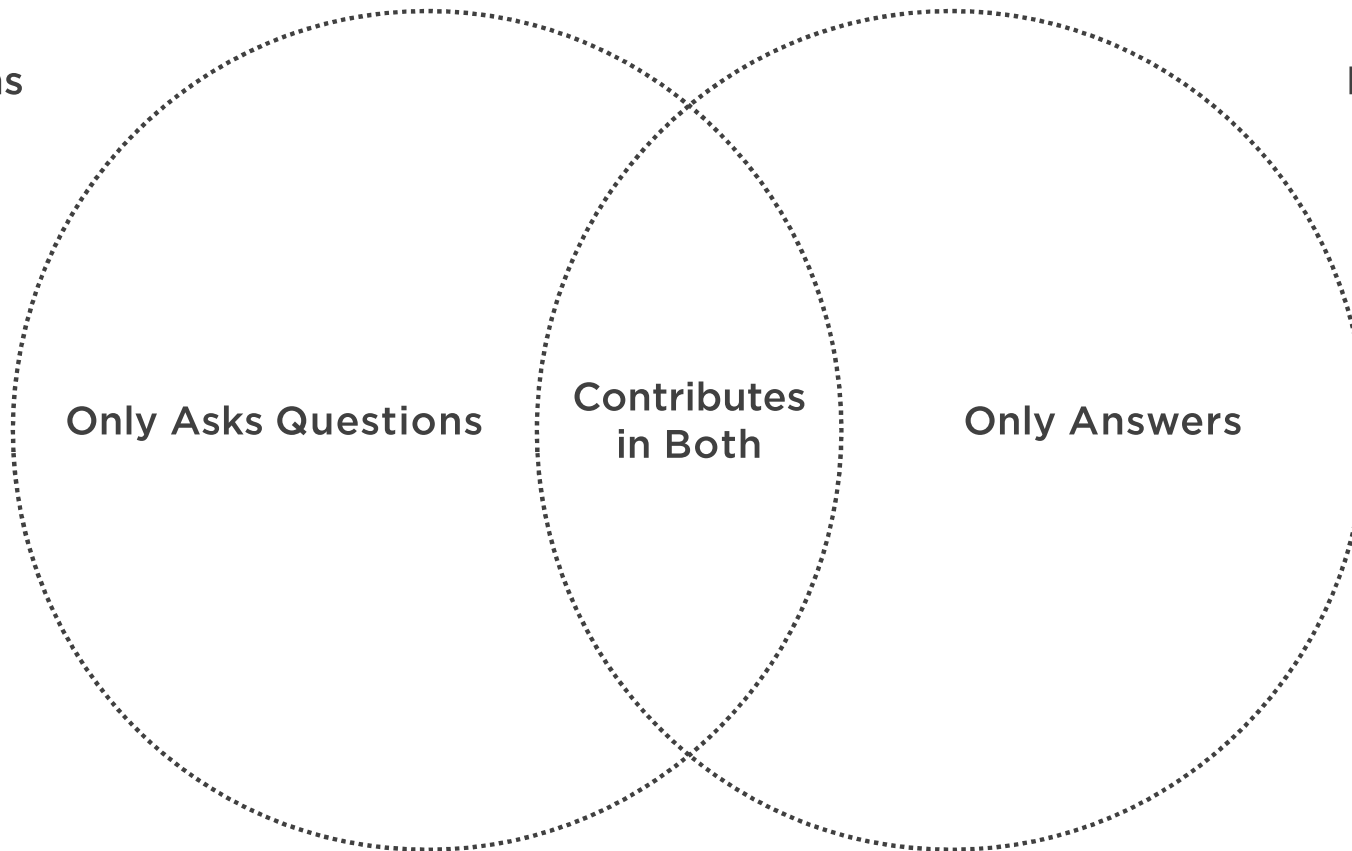
Set Operations

Posts Questions

(xavier, 1)

(troy, 2)

(xavier, 5)



Posts Answers

(xavier, 3)

(beth, 4)

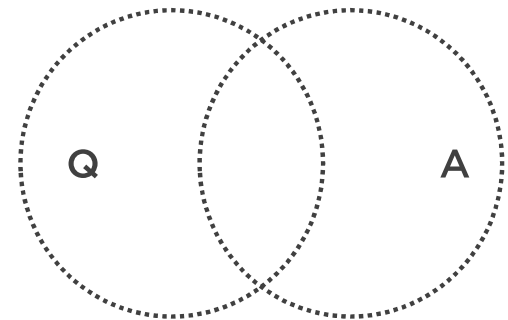


```
val questions =  
sc.parallelize(Array(("xavier",1),("troy",2),("xavier",5)))  
  
val answers =  
sc.parallelize(Array(("xavier",3),("beth",4)))  
  
questions.collect()  
answers.collect()
```

Our Data

Create with `parallelize`

If you feel confident, go for the full dataset



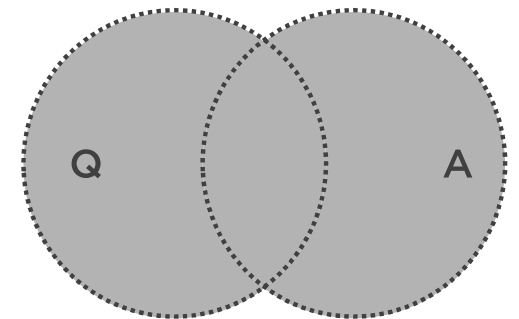
```
questions.union(answers).collect()  
questions.union(questions).collect()  
questions.union(sc  
    .parallelize(Array("irene", "juli", "luci"))).collect()
```

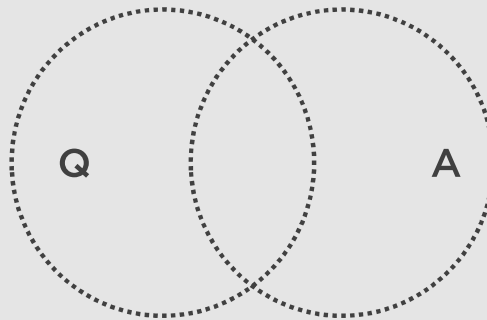
Union

RDD with all elements in both RDDs

Questions + answers

Be careful with types





(xavier, 1) (troy, 2) (xavier, 5) (xavier, 3) (beth, 4)

Union

All questions and answers

Elements remain the same



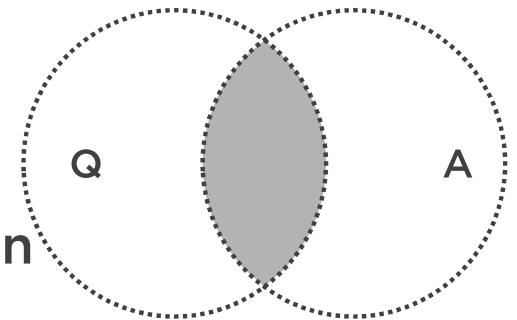

```
questions.join(answers).collect()
```

Join

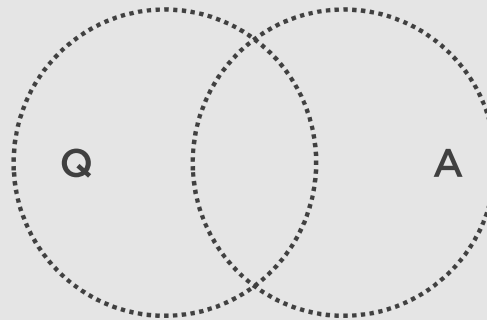
Elements with same keys in both, joined values

Hash join over the cluster, thus expensive

Unless known partitioner for narrow transformation



(troy, 2)



(beth, 4)

(xavier, (1, 3)) (xavier, (5, 3))

Join

People who have asked questions **AND** answered questions

Key is the person, value shows posts

Excludes those that do not contribute

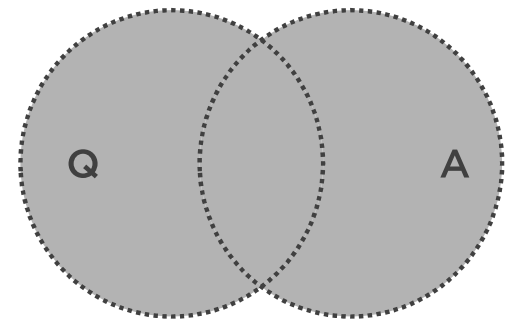


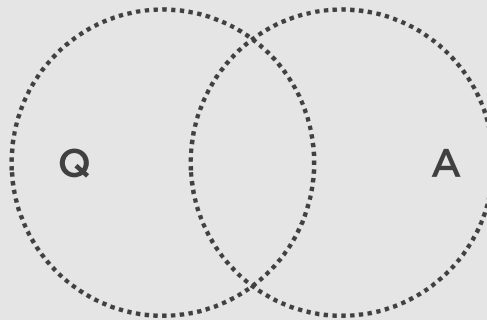
```
questions.fullOuterJoin(answers).collect()
```

fullOuterJoin

Like **join()**, but....

None where key does not appear in one RDD





(xavier, (1, **3**))

(xavier, (5, **3**))

(troy, (2, **None**))

(beth, (**None**, 4))

fullOuterJoin

All questions and answers, joined by key

- None when user does not appear in one set

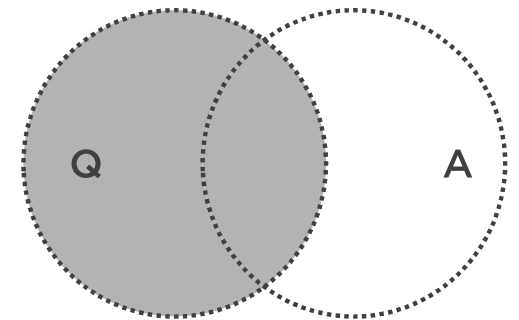


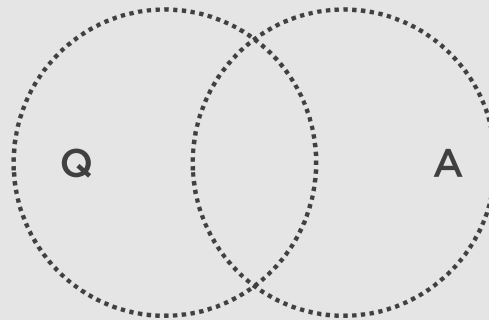
```
questions.leftOuterJoin(answers).collect()
```

leftOuterJoin

Join using keys from left set

None when key not found on right set





(beth, 4)

(xavier, (1, 3))

(xavier, (5, 3))

(troy, (2, None))

leftOuterJoin

Join on all objects from the left

List of all



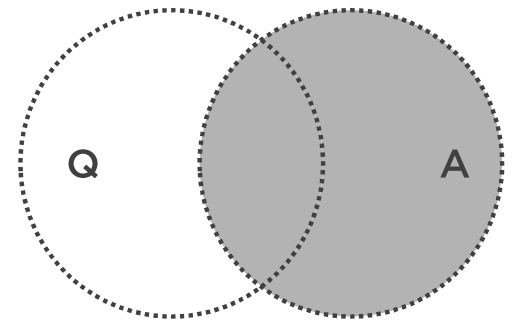
```
questions.rightOuterJoin(answers).collect()
```

rightOuterJoin

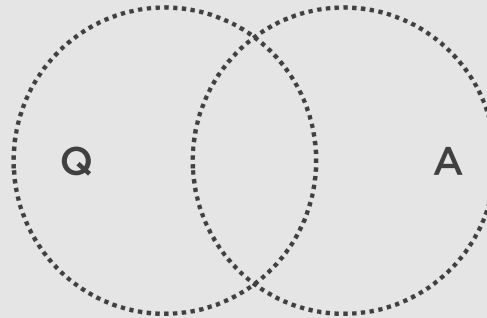
Opposite of a leftOuterJoin

Join using keys from the right set

None where keys not available in left set



(troy, 2)



(xavier, (**1** 3))

(xavier, (**5** 3))

(beth, (**None** 4))

rightOuterJoin



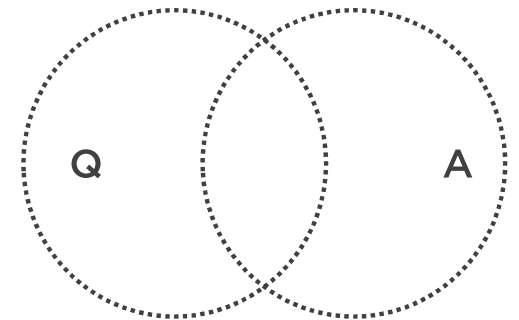

```
questions.leftOuterJoin(answers).collect()  
answers.rightOuterJoin(questions).collect()
```

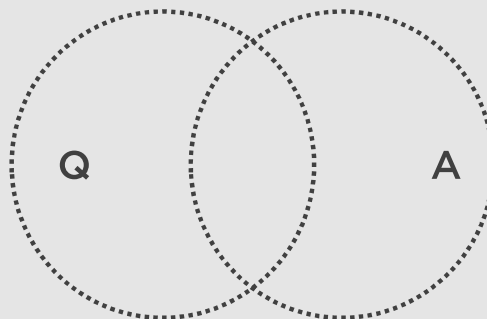
leftOuterJoin and rightOuterJoin

`questions.leftOuterJoin(answers)`

Equivalent to

`answers.rightOuterJoin(questions)`





(beth, 4)

leftOuterJoin

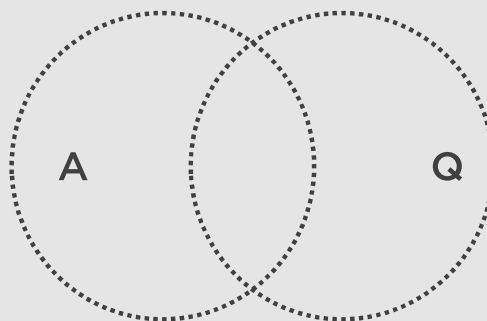
(xavier, (1, **3**))

(xavier, (5, **3**))

(troy, (2, **None**))

rightOuterJoin

(beth, 4)



(xavier, (1, **3**))

(xavier, (5, **3**))

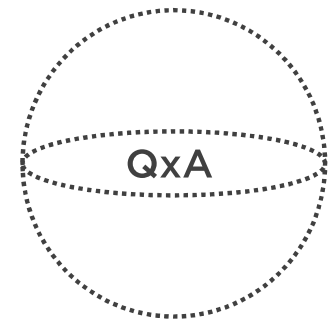
(troy, (2, **None**))

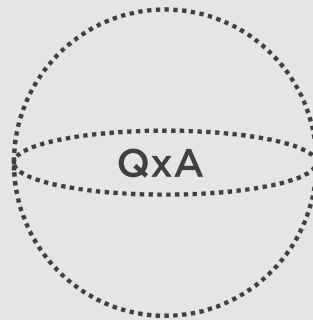
```
questions.cartesian(answers).collect()
```

Cartesian

Join of all elements in left set

With all elements in the right set





((xavier, 1) (xavier, 3)) ((xavier, 1) (beth, 4)) ((troy, 2) (xavier, 3)) ((xavier, 5) (xavier, 3)) ((troy, 2) (beth, 4)) ((xavier, 5) (beth, 4))

Cartesian



Aggregation

Grouping elements together

Foundations of Big Data analytics



```
posts_all.take(1)
val each_post_owner = posts_all.map(x => x.split(",")(6))
val posts_owner_pair_rdd = each_post_owner.map(x => (x,1))
posts_owner_pair_rdd.take(1)
```

Prepare Some Data

Extract user from each post

PairRDD

- Key is user
- Value is 1



```
val top_posters_gbk = posts_owner_pair_rdd.groupByKey()  
top_posters_gbk.take(10)
```

GroupByKey

Values grouped by each key

Data sent over the network and collected on reduce workers

Can cause problems on larger datasets



```
top_posters_gbk.map({ case (x,y) => (x, y.toList) })  
  .take(10)  
  
top_posters_gbk.map({ case (x,y) => (x, y.size) }).take(10)  
  
top_posters_gbk.map({ case (x,y) => (x, y.size) })  
  .sortBy({ case (x, y) => -y}).take(1)
```

GroupByKey

Tuple of user id and list of 1's

Posts per user? → User id and number of posts

Use sortBy for top poster




```
val top_posters_rbk = posts_owner_pair_rdd  
  .reduceByKey(_ + _)
```

ReduceByKey

Perform an operation on all elements with same key

Specify a function

Reduce operation done within partition



```
top_posters_rbk.lookup("51")
```

ReduceByKey

Perform add using `_ + _`

Pass function as parameter to `reduceByKey()`

Use `lookup()` to find top poster and confirm



```
top_posters_gbk.count()
```

```
top_posters_rbk.count()
```

groupByKey vs. reduceByKey

Do we get the same results?

Indeed we do



Preparation for aggregateByKey

```
val posts_all_entries = posts_all.map(x => x.split(", "))  
val questions = posts_all_entries.filter(x => x(1) == "1")  
val user_question_score = questions.map(x =>  
  (x(6), x(4).toInt))  
user_question_score.take(5).foreach(println)
```



```
val posts_all_entries = posts_all.map(x => x.split(","))  
val questions = posts_all_entries.filter(x => x(1) == "1")  
val user_question_score = questions.map(x =>  
    (x(6), x(4).toInt))  
user_question_score.take(5).foreach(println)
```

aggregateByKey

Like **reduceByKey()**

But takes an initial value

Specify functions for merging and combining



```
var for_keeping_count = (0,0)

def combining (tuple_sum_count: (Int, Int), next_score: Int) =
  (tuple_sum_count._1 + next_score, tuple_sum_count._2 + 1)

def merging (tuple_sum_count: (Int, Int),
tuple_next_partition_sum_count: (Int, Int)) = (tuple_sum_count._1 +
  tuple_next_partition_sum_count._1, tuple_sum_count._2 +
  tuple_next_partition_sum_count._2)

val aggregated_user_question =
  user_question_score.aggregateByKey(for_keeping_count)(combining,
    merging)

aggregated_user_question.take(1)

aggregated_user_question.lookup("51")
```



```
val aggregated_user_question =  
  user_question_score.aggregateByKey(for_keeping_count)  
  (combining, merging)
```

aggregateByKey

Combining

- Within partition

Merging

- Across partitions



```
aggregated_user_question.lookup("51")
```

aggregateByKey

Only questions, include score and user id

Define initial value, merging function, and combining function

Check with top poster




```
val user_post = questions.map(x => (x(6), x(0).toInt))
```

```
def to_list(postid: Int): List[Int] = List(postid)
```

```
def merge_posts(posta: List[Int], postb: Int) = postb ::  
posta
```

```
def combine_posts(posta: List[Int], postb: List[Int]):  
List[Int] = posta ++ postb
```



```
val combined = user_post.combineByKey(to_list, merge_posts,  
    combine_posts)  
  
combined.filter({ case (x,y) => x == "51" }).collect()
```

CombineByKey

Specify an initial value can be a function that returns a new value

Provide merge and combine functions

Like **aggregateByKey()**, but more flexible



```
user_post.lookup("51")  
user_post.countByKey()("51")
```

CountByKey

Dictionary with keys and counts of occurrences

Like a **reduceByKey()** where we count based on key



```
word_for_count.groupByKey().count()
```

```
word_for_count.reduceByKey(_ + _).count()
```

reduceByKey & groupByKey

Both can be used for the same purpose

Aggregate by keys

Work very differently underneath



Comparing groupByKey vs. reduceByKey

groupByKey



reduceByKey



groupByKey



(Spark,1)
(Spark,1)
(Spark,1)
(Spark,1)
(Spark,1)
(Spark,1)
(Spark,1)
(Spark,1)

(HUE,1)

(Cloudera,1)
(Cloudera,1)
(Cloudera,1)

12 elements vs. 5 elements

(Spark,8)

(HUE,1)

(Cloudera,3)

reduceByKey



(Spark,8)

(HUE,1)

(Cloudera,3)



Histogram



A diagram consisting of rectangles whose area is proportional to the frequency of a variable and whose width is equal to the class interval.



```
badges_reduced.take(10)
```

```
badges_reduced.map({ case (x,y) => y }).histogram(7)
```

Grouping Data into Buckets with Histogram

Histograms are very powerful graphic tools

An image is worth a thousand words

Getting the data is usually the hardest part




```
val intervals: Array[Double] =  
    Array(0,1000,2000,3000,4000,5000,6000,7000)  
  
badges_reduced.map({ case (x,y) => y }).histogram(intervals)  
  
badges_reduced.sortBy(x => -x._2).take(10)  
  
badges_reduced.filter(x => x._2 < 1000).count()
```

Grouping Data into Buckets with Histogram

Specify number of intervals

- Returns array with intervals and array of counts within intervals

Explicitly state which intervals to use



Cache

Store data for future use, to improve response times

Persist to disk, memory or both



```
reduced.setName( 'Reduced RDD' )  
reduced.cache()
```

Cache & Persist

Spark may perform caching of intermediate results

- On expensive operations, to avoid recomputing when nodes fail



Details for Job 95

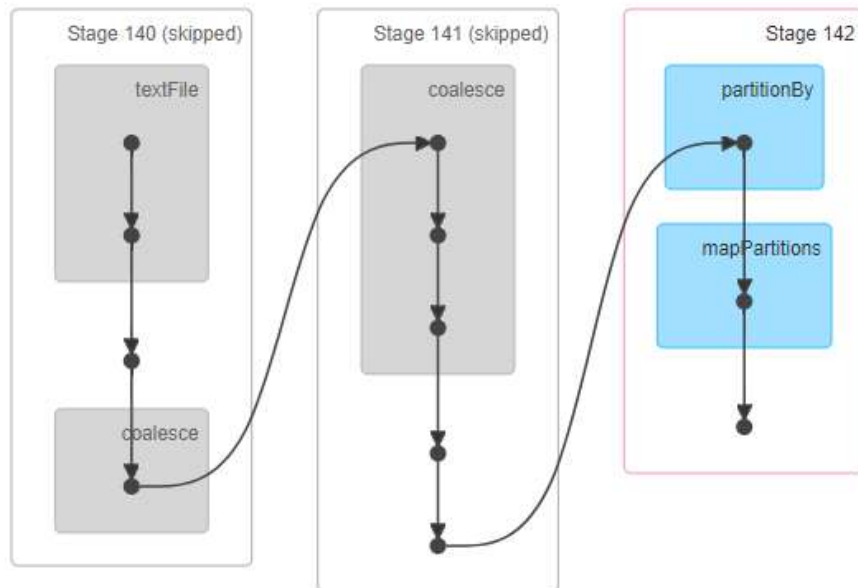
Status: SUCCEEDED

Completed Stages: 1

Skipped Stages: 2

▶ [Event Timeline](#)

▼ [DAG Visualization](#)



Completed Stages (1)

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
142	runJob at PythonRDD.scala:446 +details	2018/01/12 13:00:58	75 ms	1/1			177.8 KB	

```
reduced.cache()
```

Cache & Persist

Spark may perform caching of intermediate results

- On expensive operations, to avoid recomputing when nodes fail

If the same job called twice, entire operation may be recomputed



```
import org.apache.spark.storage.StorageLevel  
grouped.persist(StorageLevel.DISK_ONLY)
```

Cache & Persist

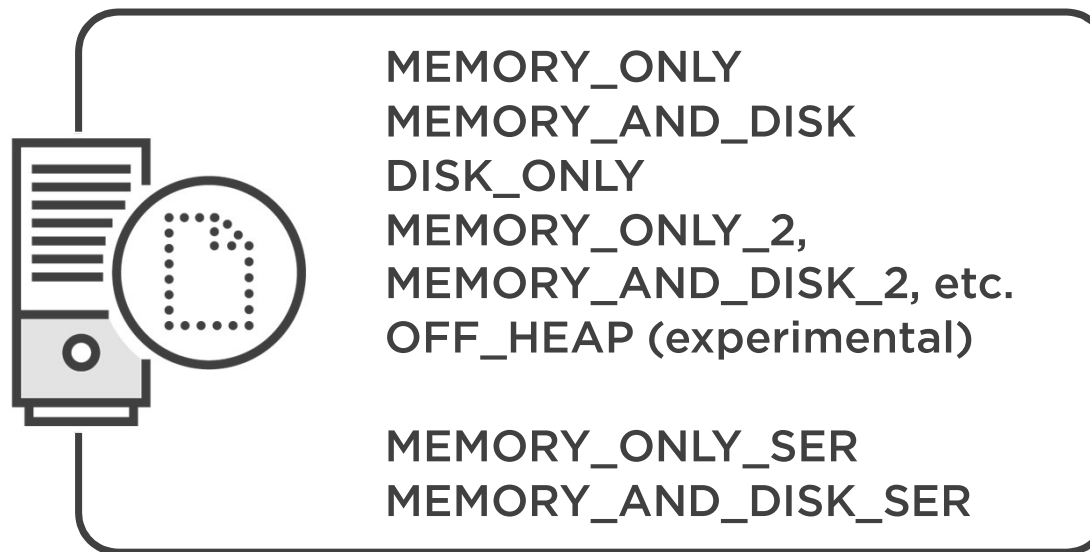
Call explicitly **cache()** and **persist()** when beneficial

- **cache()** is equivalent to **persist(MEMORY_ONLY)**

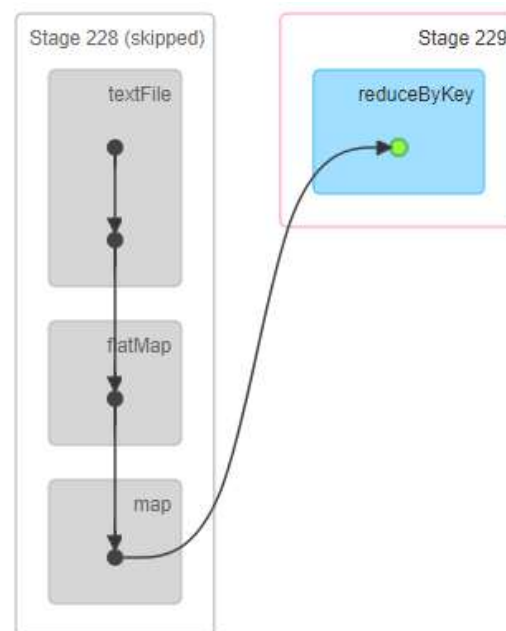
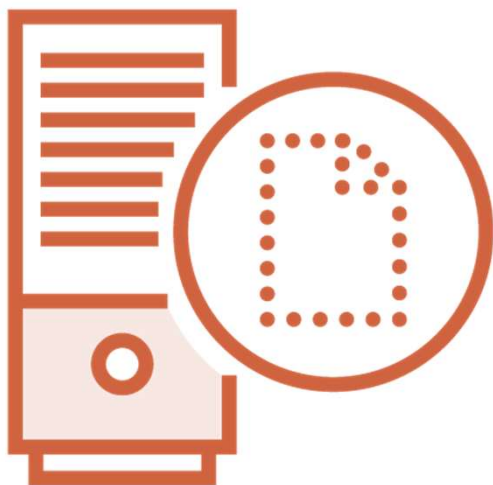
When RDD not needed anymore, call **unpersist()**



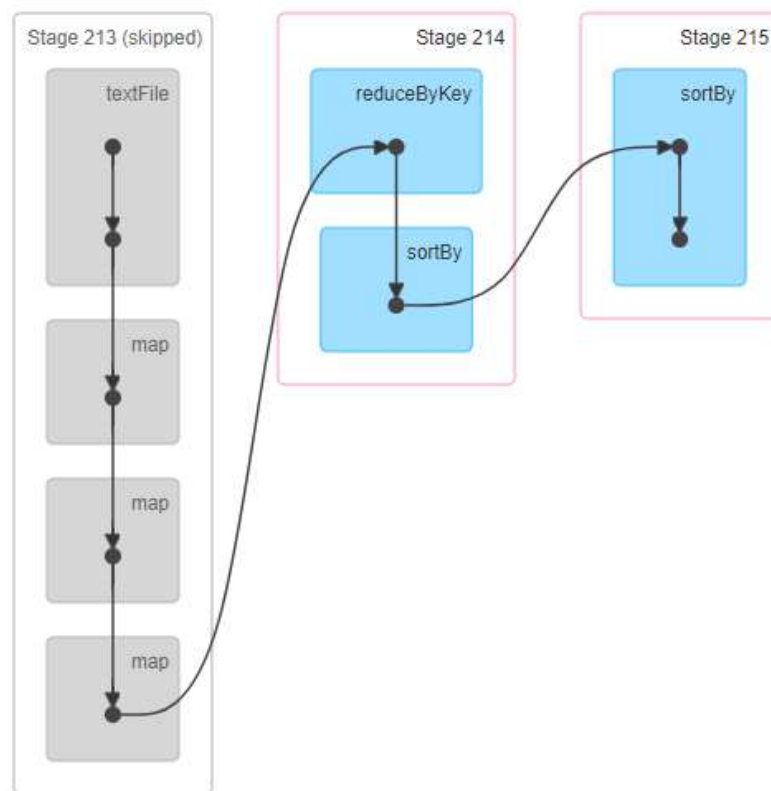
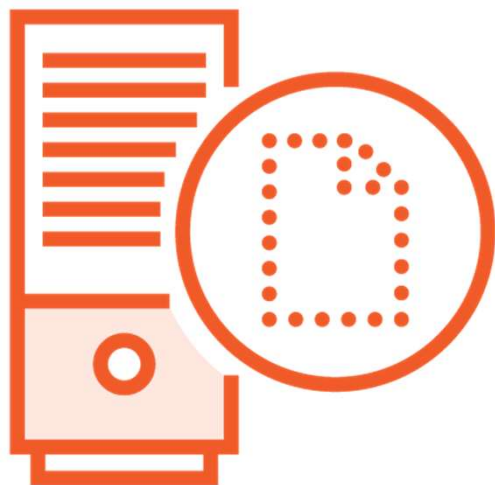
Storage Levels



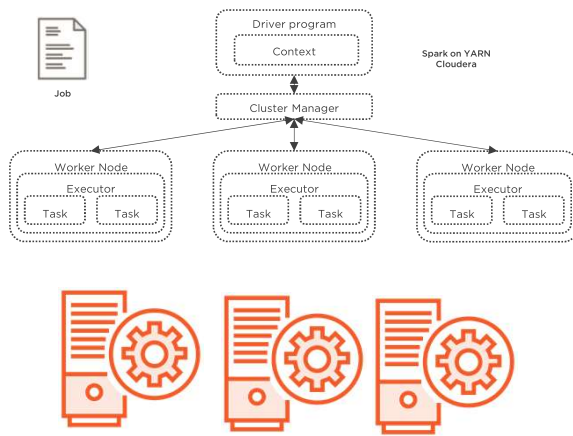
Cache & Persist



Cache in Spark UI



Spark Processing



Distributed and parallel processing

Each executor has separate copies

- Variables and functions

No propagation data back to driver

- Except on certain necessary cases
- Accumulators and broadcast variables



Shared Variables

Accumulators

"Added"

Associate and commutative

Numeric accumulator

Other types possible

Counter is one common scenario

Accumulator may not be reliable

Case of failed task

Potential duplicate counts

Broadcast Variables

Read only variable

Immutable

Fits in memory

Distributed efficiently to the cluster

Do not modify after shipped

Good case is a lookup table



```
val accumulator_badge = sc.  
  longAccumulator("Badge Accumulator")  
  
def add_badge(item: (String, String)) =  
  accumulator_badge.add(1)  
  
badges_by_badge.foreach(add_badge)  
  
accumulator_badge.value
```

Accumulator

Create accumulator and check current value

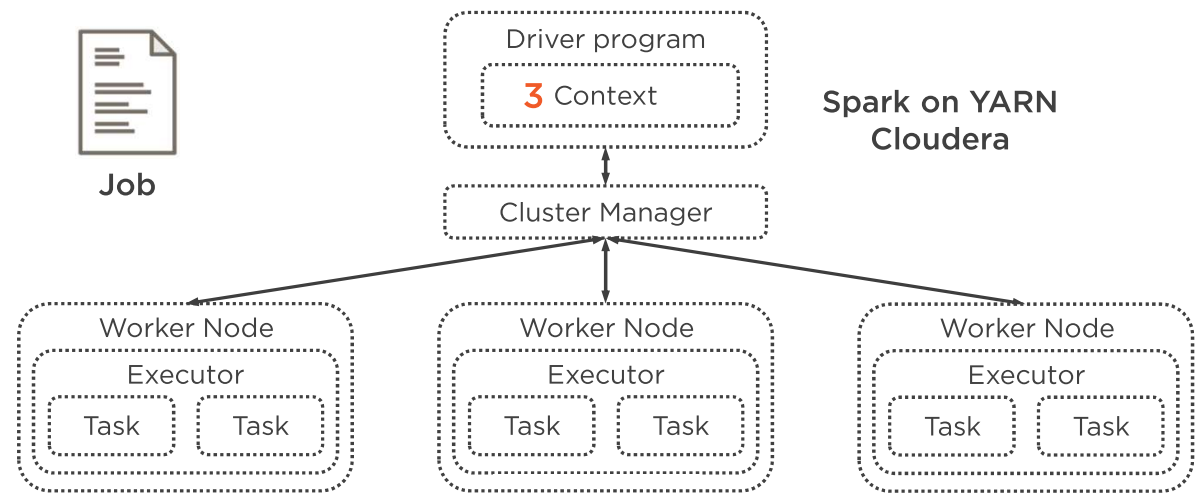
Increment accumulator function and run

Get value



Accumulators

Executors write to
accumulator in
Driver program



Aggregated Metrics by Executor

Executor ID		Address	Task Time	Total Tasks	Failed Tasks	Killed Tasks	Succeeded Tasks	Shuffle Read Size / Records	Blacklisted
88	stdout stderr	dn04.cloudera:36647	2 s	50	0	0	50	710.5 KB / 20586	0

Accumulators

Accumulable	Value
Badge Accumulator	20586

Tasks (50)

Index	ID	Attempt	Status	Locality Level	Executor ID / Host	Launch Time	Duration	GC Time	Accumulators	Shuffle Read Size / Records	Errors
0	1062	0	SUCCESS	PROCESS_LOCAL	88 / dn04.cloudera stdout stderr	2018/03/05 23:00:08	2 ms			0.0 B / 0	
1	1024	0	SUCCESS	NODE_LOCAL	88 / dn04.cloudera stdout stderr	2018/03/05 23:00:06	31 ms		Badge Accumulator: 166	9.8 KB / 166	
2	1063	0	SUCCESS	PROCESS_LOCAL	88 / dn04.cloudera stdout stderr	2018/03/05 23:00:08	2 ms			0.0 B / 0	
3	1064	0	SUCCESS	PROCESS_LOCAL	88 / dn04.cloudera stdout stderr	2018/03/05 23:00:08	2 ms			0.0 B / 0	
4	1065	0	SUCCESS	PROCESS_LOCAL	88 / dn04.cloudera stdout	2018/03/05 23:00:08	3 ms			0.0 B / 0	

```
cd users
```

```
sbt package
```

```
spark2-submit --class "PrepareUsersApp"  
                target/scala-2.11/users-project_2.11-1.0.jar
```

Convert Users.xml to CSV



Data preparation step



```
val users_all =  
sc.textFile("/user/cloudera/stackexchange/users_csv")  
users_all.take(10)  
val users_columns = users_all.map(split_the_line)  
users_columns.take(3)  
  
top_posters_rbk.take(10)  
top_posters_rbk.lookup("51")
```




```
def get_name(user_column: Array[String]) = {  
    val user_id = user_column(0)  
    val user_name = user_column(3)  
    var user_post_count = "0"  
    if (broadcast_tp.value.keySet.exists(_ == user_id))  
        user_post_count = broadcast_tp.value(user_id).toString  
    (user_id, user_name, user_post_count)  
}
```



```
val tp = top_posters_rbk.collectAsMap()  
val broadcast_tp = sc.broadcast(tp)
```

Broadcast Variable

Create a broadcast variable using the context

Access when necessary, i.e. lookup

Use **value**



```
val user_info = users_columns.map(get_name)  
user_info.take(10)
```

Broadcast Variable

Create using **sc.broadcast()**

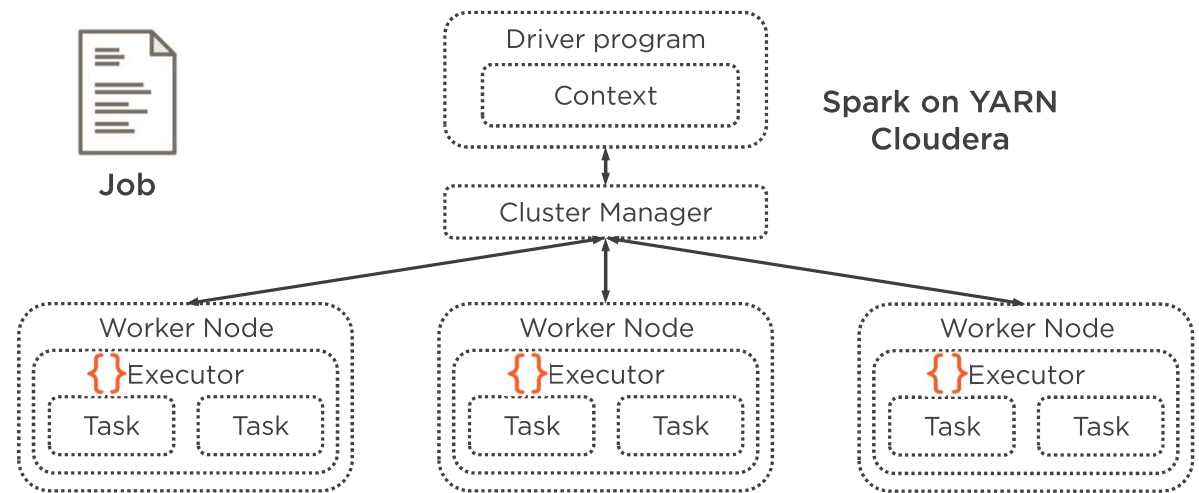
- Assign to a **variable**

Access using **variable.value**



Broadcast Variables

Executors read
from Broadcast
variable



Developing Self-contained Spark Apps



Requires

- Create the SparkContext
- Dependencies
- Execute using spark2-submit



```
import org.apache.spark.{SparkConf, SparkContext}
val conf = new SparkConf()
    .setMaster("yarn")
    .setAppName("Self Contained Application")
sc = new SparkContext(conf)
```

Creating the SparkContext

Corresponding **import**

Create **sc**



```
cd users
```

```
sbt package
```

Bundling Your Application

Use **sbt** or **Maven** to create an "uber jar"

Use **sbt package**, respect folder structure

We covered in the Scala refresher module



```
spark2-submit --class "PrepareUsersApp"  
               target/scala-2.11/users-project_2.11-1.0.jar
```

Launching Application

Using **spark2-submit**

Code to be executed, submitted as a job




```
spark2-submit --class "PrepareUsersApp"  
              --jars my_dependency.jar  
              --packages com.databricks:spark-xml_2.11:0.4.1  
              target/scala-2.11/users-project_2.11-1.0.jar
```

Dependencies

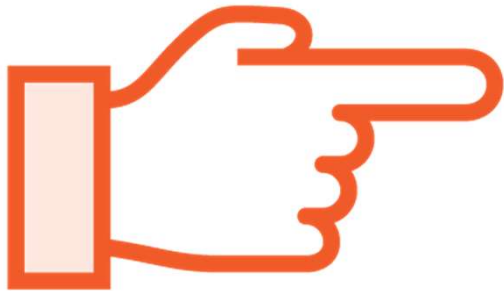
Use **jars** parameter

- Supports file, hdfs, http, ftp or local, but no directory expansion

Maven coordinates with **packages**



Disadvantages of RDDs



Don't take this the wrong way

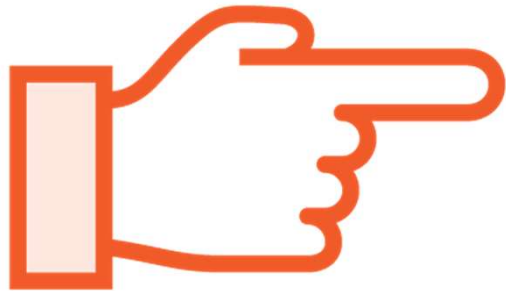
RDDs are still used, even internally

Extremely powerful

Limitations on potential optimizations



Disadvantages of RDDs



Performance

Schema less

Steeper learning curve

"Everybody knows SQL"





Takeaway



Anonymous Functions

- Lambdas

Transformations vs. Actions

- Transformations return RDDs
- Actions trigger computation



Takeaway



Map, FlatMap, Filter, Sort, ...

Partitions

Sampling

Set operations

Aggregations



Takeaway



Histogram

Caching & Persisting

Shared variables

Self contained applications



Takeaway



Disadvantages of RDDs

