

Understanding a Typed API: Datasets



Xavier Morera

HELPING DEVELOPERS UNDERSTAND SEARCH & BIG DATA

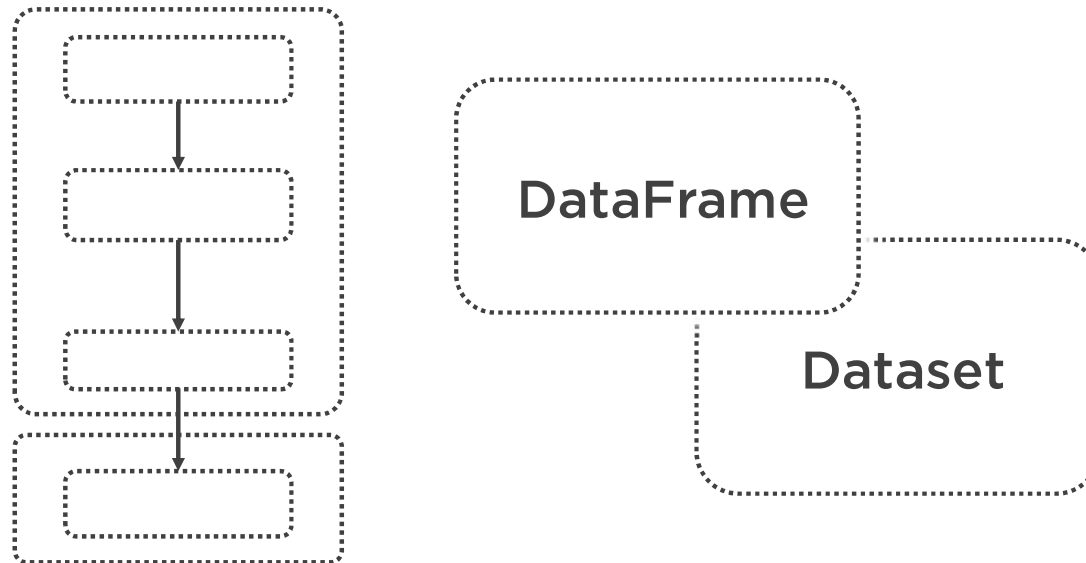
@xmorera www.xaviermorera.com



Understanding a Typed API: Datasets



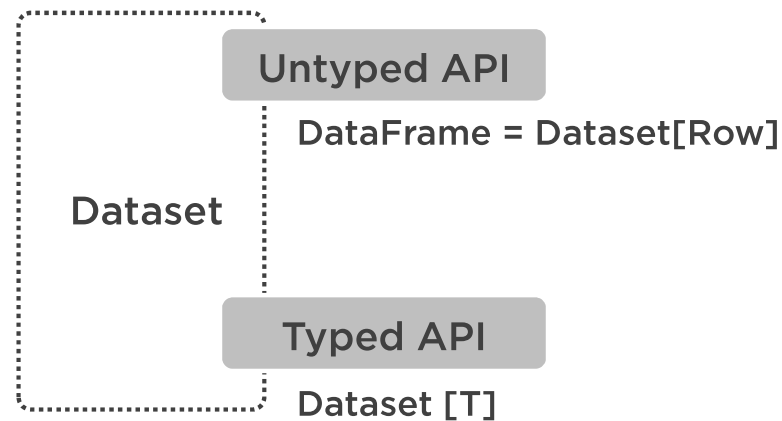
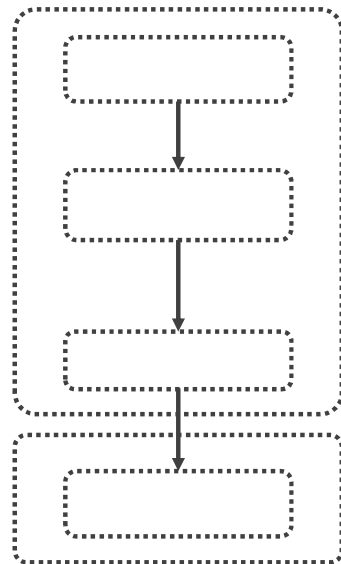
Understanding a Typed API: Datasets



Since 1.x



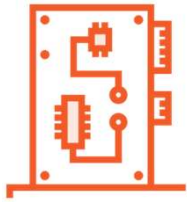
Understanding a Typed API: Datasets



Since 2.0



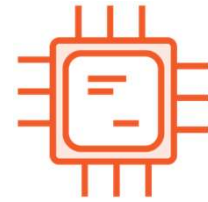
The Motivation Behind Datasets



Network



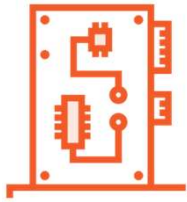
Storage



CPU



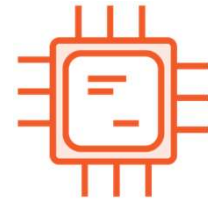
The Motivation Behind Datasets



Network
1 GBPS



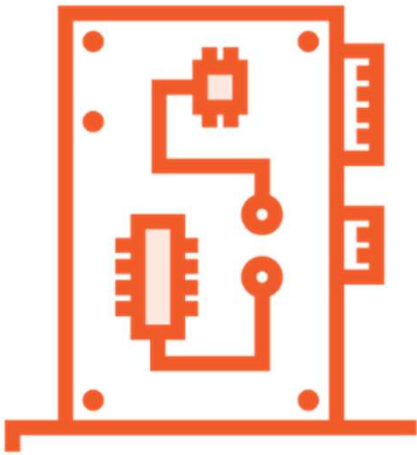
Storage



CPU



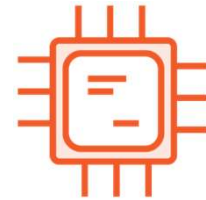
The Motivation Behind Datasets



Network
10 Gbps



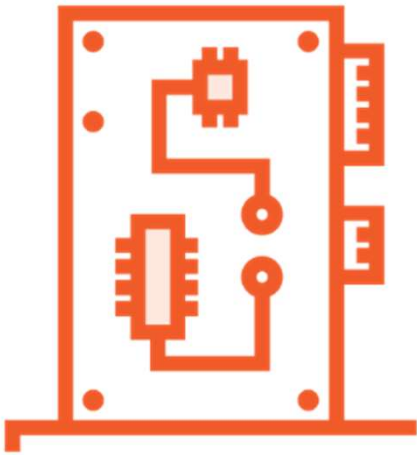
Storage



CPU



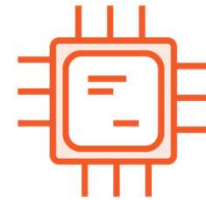
The Motivation Behind Datasets



Network
10 Gbps



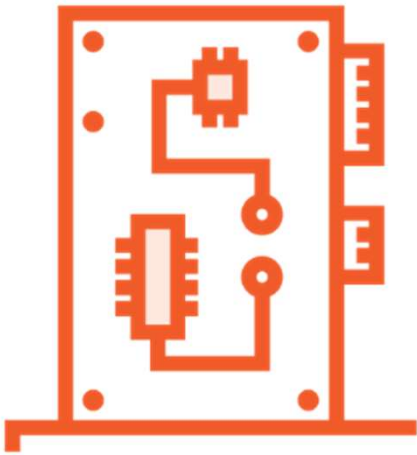
Storage
50 Mbps



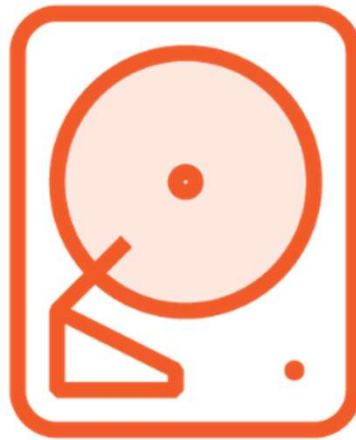
CPU



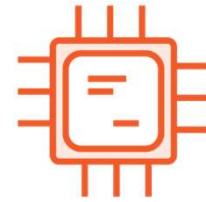
The Motivation Behind Datasets



Network
10 Gbps



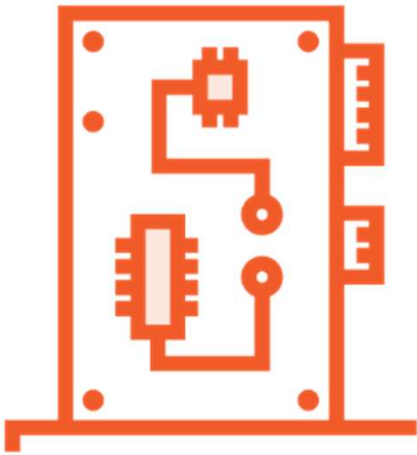
Storage
500 Mbps



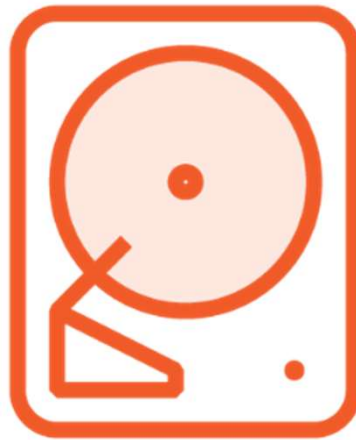
CPU
3 Ghz



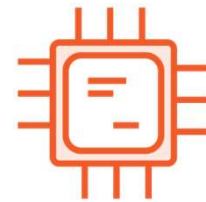
The Motivation Behind Datasets



Network
10 Gbps



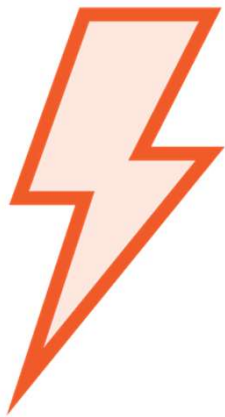
Storage
500 Mbps



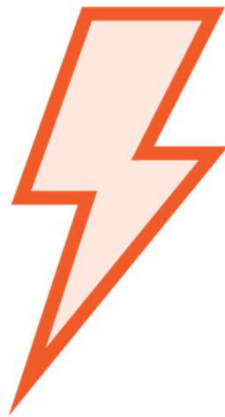
CPU
3 Ghz



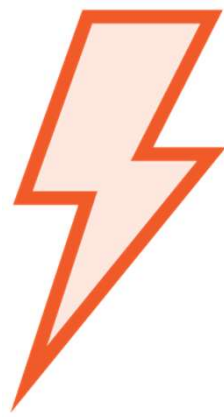
The Motivation Behind Datasets



Network
10 Gbps



Storage
500 Mbps



CPU
3 Ghz



```
val postsRDD =  
sc.textFile("/user/cloudera/stackexchange/simple_titles.txt")  
  
val postsDS =  
spark.read.text("/user/cloudera/stackexchange/simple_titles.txt").  
  as[String]  
postsRDD.setName("postsRDD")  
  
val postsDF =  
spark.read.text("/user/cloudera/stackexchange/simple_titles.txt")
```

Size Matters







Create an RDD and Dataset

- Cache and count
- Same with a DataFrame

Think beyond MBs, imagine TB or PB




```
spark.sql("Selct Score from PostsSE").show(5)
spark.sql("Select Scre from PostsSE").show(5)
postsDF.selct("Score").show(5)
postsDF.select($"Scre").first
postsDS.first.Scre
val firstScore: String = postsDS.first.Score
```

	SQL	DataFrames	Datasets
Syntax errors			
Analysis errors			



What's a Dataset

"""	#	

Dataset

- Strongly typed collection of objects
- Domain specific objects
- Type safe
- Functional programming
- Query optimization

You know the API

- More than just Row



What Can You Store in a Dataset?

.as[Post]
.toDS()

Encoder

- JVM objects <-> Spark's internal representation

Supported objects

- Primitive types
- Complex types
- Product objects
- Row objects



DataFrames

Remember Row objects?



DataFrame - org.apache

← → ↻ 🔒 Not secure | spark.apache.org/docs/1.6.3/api/scala/index.html#org.apache.spark.sql.DataFrame ☆


dataframe

#ABCDEFGHIJKLMNOPQRSTUVWXYZ

display packages only

org.apache.spark.sqlhide focus

- DataFrame
- DataFrameHolder
- DataFrameNaFunctions
- DataFrameReader
- DataFrameStatFunctions
- DataFrameWriter

 org.apache.spark.sql

DataFrame

class **DataFrame** extends Queryable with Serializable

A distributed collection of data organized into named columns.

A [DataFrame](#) is equivalent to a relational table in Spark SQL. The following example creates a [DataFrame](#) by pointing Spark SQL to a Parquet data set.

```
val people = sqlContext.read.parquet("...") // in Scala
DataFrame people = sqlContext.read().parquet("...") // in Java
```

Once created, it can be manipulated using the various domain-specific-language (DSL) functions defined in: [DataFrame](#) (this class), [Column](#), and [functions](#).

To select a column from the data frame, use apply method in Scala and col in Java.

```
val ageCol = people("age") // in Scala
Column ageCol = people.col("age") // in Java
```

Note that the [Column](#) type can also be manipulated through its various functions.

```
// The following creates a new column that increases everybody's age by 10.
people("age") + 10 // in Scala
people.col("age").plus(10); // in Java
```

A more concrete example in Scala:

```
// To create DataFrame using SQLContext
val people = sqlContext.read.parquet("...")
val department = sqlContext.read.parquet("...")

people.filter("age > 30")
  .join(department, people("deptId") === department("id"))
  .groupBy(department("name"), "gender")
  .agg(avg(people("salary")), max(people("age")))
```

and in Java:

```
// To create DataFrame using SQLContext
DataFrame people = sqlContext.read().parquet("...");
DataFrame department = sqlContext.read().parquet("...");
```

Experimental

Dataset - org.apache.spark

← → ↻ 🔒 Not secure | spark.apache.org/docs/1.6.3/api/scala/index.html#org.apache.spark.sql.Dataset

dataset

#ABCDEFGHIJKLMNOPQRSTUVWXYZ

display packages only

org.apache.spark.sql

hide focus

Dataset

DatasetHolder

GroupedDataset

org.apache.spark.sql

Dataset

class **Dataset**[T] extends Queryable with Serializable with Logging

A [Dataset](#) is a strongly typed collection of objects that can be transformed in parallel using functional or relational operations.

Experimental

A [Dataset](#) differs from an RDD in the following ways:

- Internally, a [Dataset](#) is represented by a Catalyst logical plan and the data is stored in the encoded form. This representation allows for additional logical operations and enables many operations (sorting, shuffling, etc.) to be performed without deserializing to an object.
- The creation of a [Dataset](#) requires the presence of an explicit [Encoder](#) that can be used to serialize the object into a binary format. Encoders are also capable of mapping the schema of a given object to the Spark SQL type system. In contrast, RDDs rely on runtime reflection based serialization. Operations that change the type of object stored in the dataset also need an encoder for the new type.

A [Dataset](#) can be thought of as a specialized [DataFrame](#), where the elements map to a specific JVM object type, instead of to a generic [Row](#) container. A [DataFrame](#) can be transformed into specific [Dataset](#) by calling `df.as[ElementType]`. Similarly you can transform a strongly-typed [Dataset](#) to a generic [DataFrame](#) by calling `ds.toDF()`.

COMPATIBILITY NOTE: Long term we plan to make [DataFrame](#) extend `Dataset[Row]`. However, making this change to the class hierarchy would break the function signatures for the existing functional operations (map, flatMap, etc). As such, this class should be considered a preview of the final API. Changes will be made to the interface after Spark 1.6.

Annotations

@Experimental()

Source

[Dataset.scala](#)

Since

1.6.0

► Linear Supertypes

q

Ordering

Grouped

Alphabetic

By inheritance

Inherited

Dataset

Logging

Serializable

Serializable

Queryable

AnyRef

Any

Hide All

Show all

Learn more about member selection

Visibility

Public

All

basic

Dataset

#ABCDEFGHIJKLMNOPQRSTUVWXYZ

— deprecated

display packages only

org.apache.spark.sql hide focus

- Dataset
- DatasetHolder
- KeyValueGroupedDataset
- RelationalGroupedDataset



org.apache.spark.sql

Dataset

Related Doc: [package sql](#)

class **Dataset**[T] extends Serializable

A Dataset is a strongly typed collection of domain-specific objects that can be transformed in parallel using functional or relational operations. Each Dataset also has an untyped view called a DataFrame, which is a Dataset of [Row](#).

Operations available on Datasets are divided into transformations and actions. Transformations are the ones that produce new Datasets, and actions are the ones that trigger computation and return results. Example transformations include map, filter, select, and aggregate (groupBy). Example actions count, show, or writing data out to file systems.

Datasets are "lazy", i.e. computations are only triggered when an action is invoked. Internally, a Dataset represents a logical plan that describes the computation required to produce the data. When an action is invoked, Spark's query optimizer optimizes the logical plan and generates a physical plan for efficient execution in a parallel and distributed manner. To explore the logical plan as well as optimized physical plan, use the explain function.

To efficiently support domain-specific objects, an [Encoder](#) is required. The encoder maps the domain specific type T to Spark's internal type system. For example, given a class Person with two fields, name (string) and age (int), an encoder is used to tell Spark to generate code at runtime to serialize the Person object into a binary structure. This binary structure often has much lower memory footprint as well as are optimized for efficiency in data processing (e.g. in a columnar format). To understand the internal binary representation for data, use the schema function.

There are typically two ways to create a Dataset. The most common way is by pointing Spark to some files on storage systems, using the read function available on a SparkSession.

```
val people = spark.read.parquet("...").as[Person] // Scala
Dataset<Person> people = spark.read().parquet("...").as(Encoders.bean(Person.class)); // Java
```

Datasets can also be created through transformations available on existing Datasets. For example, the following creates a new Dataset by applying a filter on the existing one:

```
val names = people.map(_.name) // in Scala; names is a Dataset[String]
Dataset<String> names = people.map((Person p) -> p.name, Encoders.STRING);
```

Dataset operations can also be untyped, through various domain-specific-language (DSL) functions defined in: Dataset (this

DataFrame

#ABCDEFGHIJKLMNOPQRSTUVWXYZ

— deprecated

display packages only

org.apache.spark.sql hide focus

- DataFrameNaFunctions
- DataFrameReader
- DataFrameStatFunctions
- DataFrameWriter



package spark

Core Spark functionality. [org.apache.spark.SparkContext](#) serves as the main entry point to Spark, while [org.apache.spark.rdd.RDD](#) is the data type representing a distributed collection, and provides most parallel operations.

In addition, [org.apache.spark.rdd.PairRDDFunctions](#) contains operations available only on RDDs of key-value pairs, such as `groupByKey` and `join`; [org.apache.spark.rdd.DoubleRDDFunctions](#) contains operations available only on RDDs of Doubles; and [org.apache.spark.rdd.SequenceFileRDDFunctions](#) contains operations available on RDDs that can be saved as `SequenceFiles`. These operations are automatically available on any RDD of the right type (e.g. `RDD[(Int, Int)]`) through implicit conversions.

Java programmers should reference the [org.apache.spark.api.java](#) package for Spark programming APIs in Java.

Classes and methods marked with **Experimental** are user-facing features which have not been officially adopted by the Spark project. These are subject to change or removal in minor releases.

Classes and methods marked with **Developer API** are intended for advanced users want to extend Spark through lower level interfaces. These are subject to changes or removal in minor releases.

Source [package.scala](#)

► Linear Supertypes



Ordering **Alphabetic** By Inheritance

Inherited **spark** AnyRef Any

Hide All **Show All**

Visibility **Public** All

DISK_ONLY_2
[StorageLevel](#)

DOUBLE
[Encoders](#)

DStream
[dstream](#)

DataFrame
[sql](#)

DataFrameNaFunctions
[sql](#)

DataFrameReader
[sql](#)

DataFrameStatFunctions
[sql](#)

DataFrameWriter
[sql](#)

Type Members

- ▶ class **AnalysisException** extends Exception with Serializable
Thrown when a query fails to analyze, usually because the query itself is invalid.

- ▶ class **Column** extends [Logging](#)
A column that will be computed based on the data in a DataFrame.

- ▶ class **ColumnName** extends [Column](#)
A convenient class used for constructing schema.

- type **DataFrame** = [Dataset](#)[[Row](#)]

- ▶ final class **DataFrameNaFunctions** extends AnyRef
Functionality for working with missing data in DataFrames.

- ▶ class **DataFrameReader** extends [Logging](#)
Interface used to load a [Dataset](#) from external storage systems (e.g.

- ▶ final class **DataFrameStatFunctions** extends AnyRef

We already know quite
about Datasets!



Dataset[Row]



Dataset[T]



~~new~~
apply()

Case Class

Like a regular Scala class

- But with a few differences

Modeling immutable data

An instance is called a Product

Initialized a little bit differently



```
case class Post(Id: Integer, UserId: String, Score: Integer)
val post = Post(1, "1", 25)
post.Id
post.UserId
post.UserId = "2"
```

Case Class

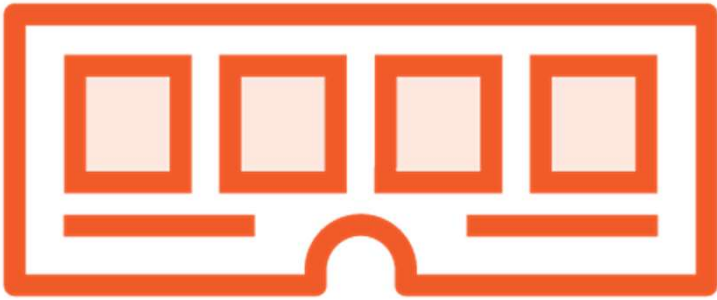
Define **case class**

On object construction **new** is not required

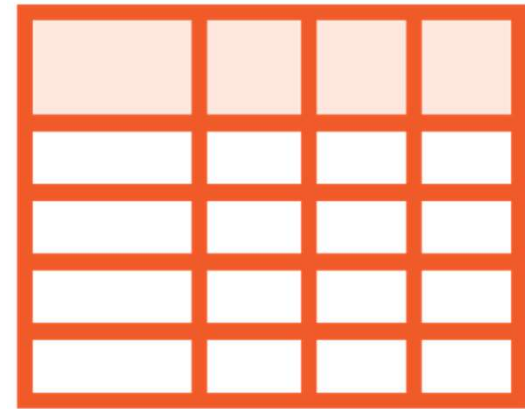
- By default **apply** method is used



Creating Datasets



Memory
`.toDS()`



DataFrame
`.as[CaseClass]`



```
val primitiveDS = Seq(10, 20, 30).toDS()
val differentDS = Seq(10, "20", 30).toDS()
primitiveDS.map(_ + 1).show()
val complexDS = Seq(("Xavier", 1), ("Irene", 2)).toDS()
complexDS.map(x => x._2 + 10).show()
```

From Data in Memory

Use **createDataset()**, with **toDS()**

Q: What did we use a lot in RDDs but not in DataFrames?

- Functional programming
- Higher-order functions with typed transformations



```
val postsRDD =  
sc.textFile("/user/cloudera/stackexchange/simple_titles.txt")  
  
val postsDSfromRDD = spark.createDataset(postsRDD)  
  
postsDSfromRDD.show(5)
```

Using RDDs

Create Dataset from an RDD

Also using **createDataset()**



```
import org.apache.spark.sql.types._
val postsSchema =
  StructType(Array(
    StructField("Id", IntegerType),
    StructField("PostTypeId", IntegerType),
    StructField("AcceptedAnswerId", IntegerType),
    StructField("CreationDate", TimestampType),
    StructField("Score", IntegerType),
    StructField("ViewCount", IntegerType),
    StructField("OwnerUserId", IntegerType),
    StructField("LastEditorUserId", IntegerType),
    StructField("LastEditDate", TimestampType),
    StructField("Title", StringType),
    StructField("LastActivityDate", TimestampType),
    StructField("Tags", StringType),
    StructField("AnswerCount", IntegerType),
    StructField("CommentCount", IntegerType),
    StructField("FavoriteCount", IntegerType)))
```



```
case class Post(Id: Int, PostTypeId: Int, Score: Integer, ViewCount: Integer, AnswerCount: Integer, OwnerUserId: Integer)
```

```
val posts_all = spark.read.schema(postsSchema)  
  .csv("/user/cloudera/stackexchange/posts_all_csv")
```

```
val postsDF = posts_all.select($"Id", $"PostTypeId", $"Score",  
  $"ViewCount", $"AnswerCount", $"OwnerUserId")
```




```
val postsDSfromDF = postsDF.as[Post]
postsDSfromDF.show(5)
postsDSfromDF
  .groupByKey(row => row.OwnerUserId).count().show()
```

From DataFrames with Case Classes

Create a Dataset from a DataFrame

- Using `as[]` and the case class

Perform strongly typed operations



Provide an API that allows performing transformations just like RDDs But with the performance and robustness of the Spark SQL execution engine

Dataset API



Typed Transformations

Untyped Transformations



```
postsDS
```

```
val postsLessDS = postsDS.filter('ViewCount < 533)
```

```
postsLessDS
```

Typed Transformations

Return a Dataset

- Type information is preserved

i.e. **filter**, **sort**, **distinct**...



```
val postsNotDS = postsDS.select('Id, 'ViewCount')  
postsNotDS
```

Untyped Transformations

Return a DataFrame

- Type information not preserved

i.e. **select**, **groupBy**, **join**...



Typed Transformations

- Filter
- Distinct
- Limit

Untyped Transformations

- Select
- Join
- GroupBy



```
postsDS.describe("ViewCount").show()

val postsLessDS = postsDS.filter(p => (p.ViewCount == 533)).show()

postsDS.filter(p => p.OwnerUserId == 51).count()

val mini_postsDS = postsDS.map(d => (d.Id, d.Score))

val mini_postsDF = postsDS.select($"Id", $"Score")

postsDS.groupBy($"UserId")

postsDS.groupBy($"UserId"). //tab
```

Dataset Operations

- **Transformations with domain specific objects**
 - What we learned earlier still applies
 - Create new Datasets

Explore the Dataset API



Performance

RDD

Lower-level API

- Unstructured data
- Fine tune
- Manage low level details
- Complex data types

DataFrame

Untyped Higher-level API

- Structured data
- Semi-structured data
- "Think in SQL"
- Performance is key

Dataset

Typed Higher-level API

- Structured data
- Semi-structured data
- Type safety
- Functional APIs



Performance

RDD



Dataset



DataFrame

Lower-level API

Unstructured data
Fine tune
Manage low level details
Complex data types

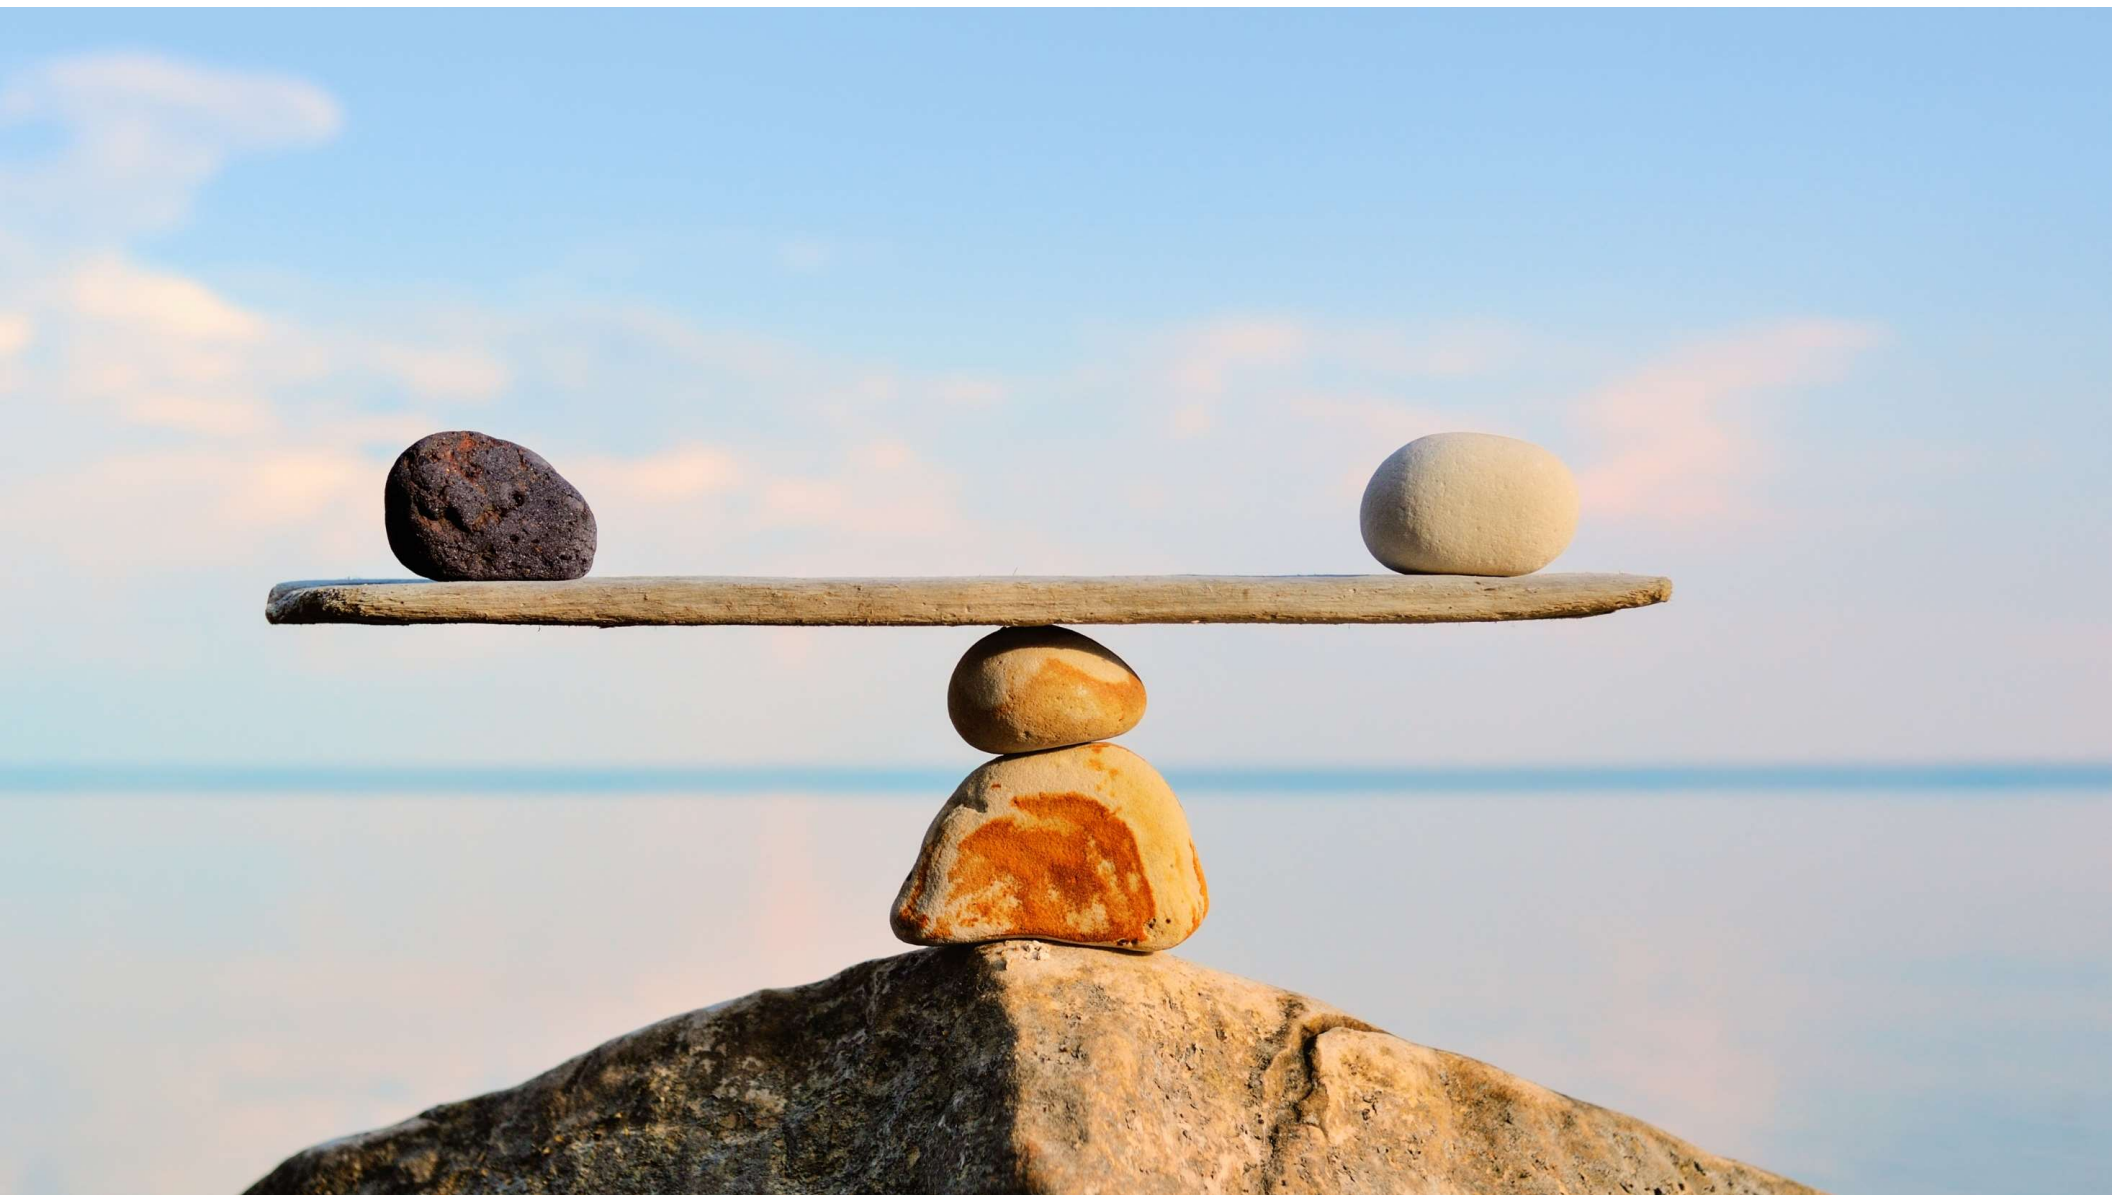
Typed Higher-level API

Structured data
Semi-structured data
Type safety
Functional APIs

Untyped Higher-level API

Structured data
Semi-structured data
"Think in SQL"
Performance is key





Takeaway



Typed API: Datasets

Motivation behind Datasets

- Performance

Syntax Errors vs. Analysis Errors



Takeaway



What's a Dataset?

- Strongly typed collection
- Domain specific objects

What can we store in a Dataset?

- Primitive types
- Complex types
- Product objects
- Row objects



Takeaway



Create Datasets

Typed transformations

Untyped transformations

Higher-order functions

Explore the API

RDDs vs. DataFrames vs. Datasets

